

# Computer Architecture: A Short Introduction

Alexandre Bartel

2019

# What is a computer?

- ▶ Input devices (keyboard, mouse, ...)
- ▶ Processing units (CPU, ...)
- ▶ Output devices (screen, ...)



image source: openclipart.com

# Inside a Computer

1. Screen
2. Motherboard
3. Central Processing Unit (CPU)
4. Memory
5. Extension cards
6. Power Supply
7. Bluray reader
8. Hard drive
9. Keyboard
10. Mouse

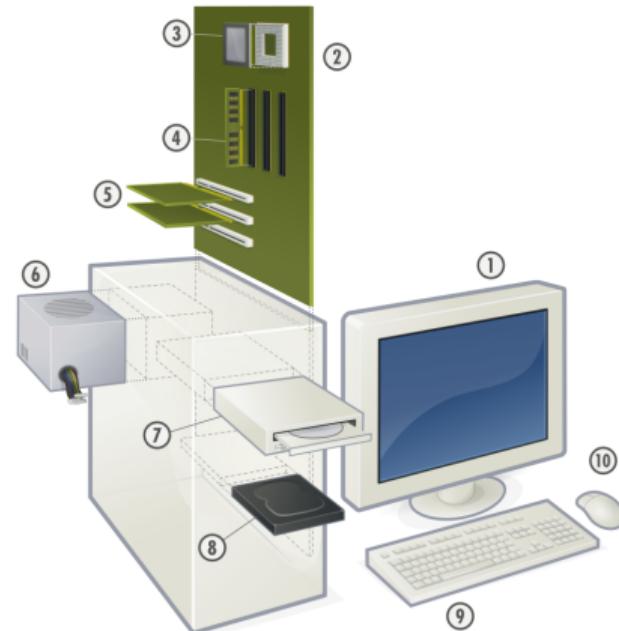


image source: wikimedia.org

# Components Identification

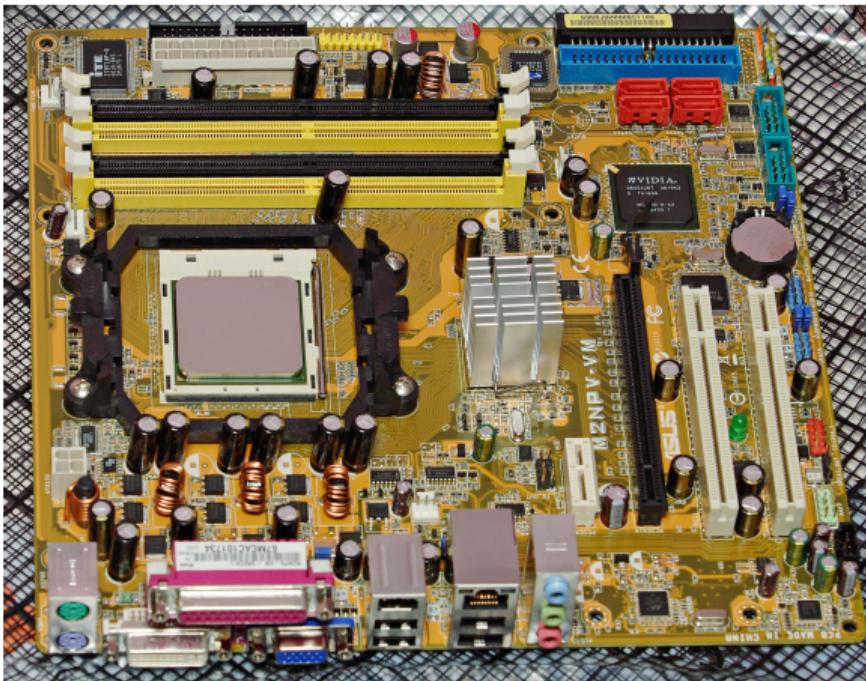


image source: wikipedia.org

# Components Identification

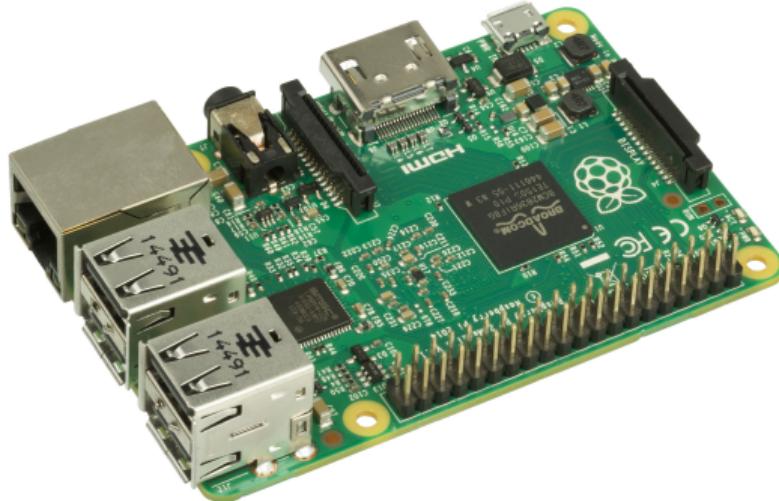
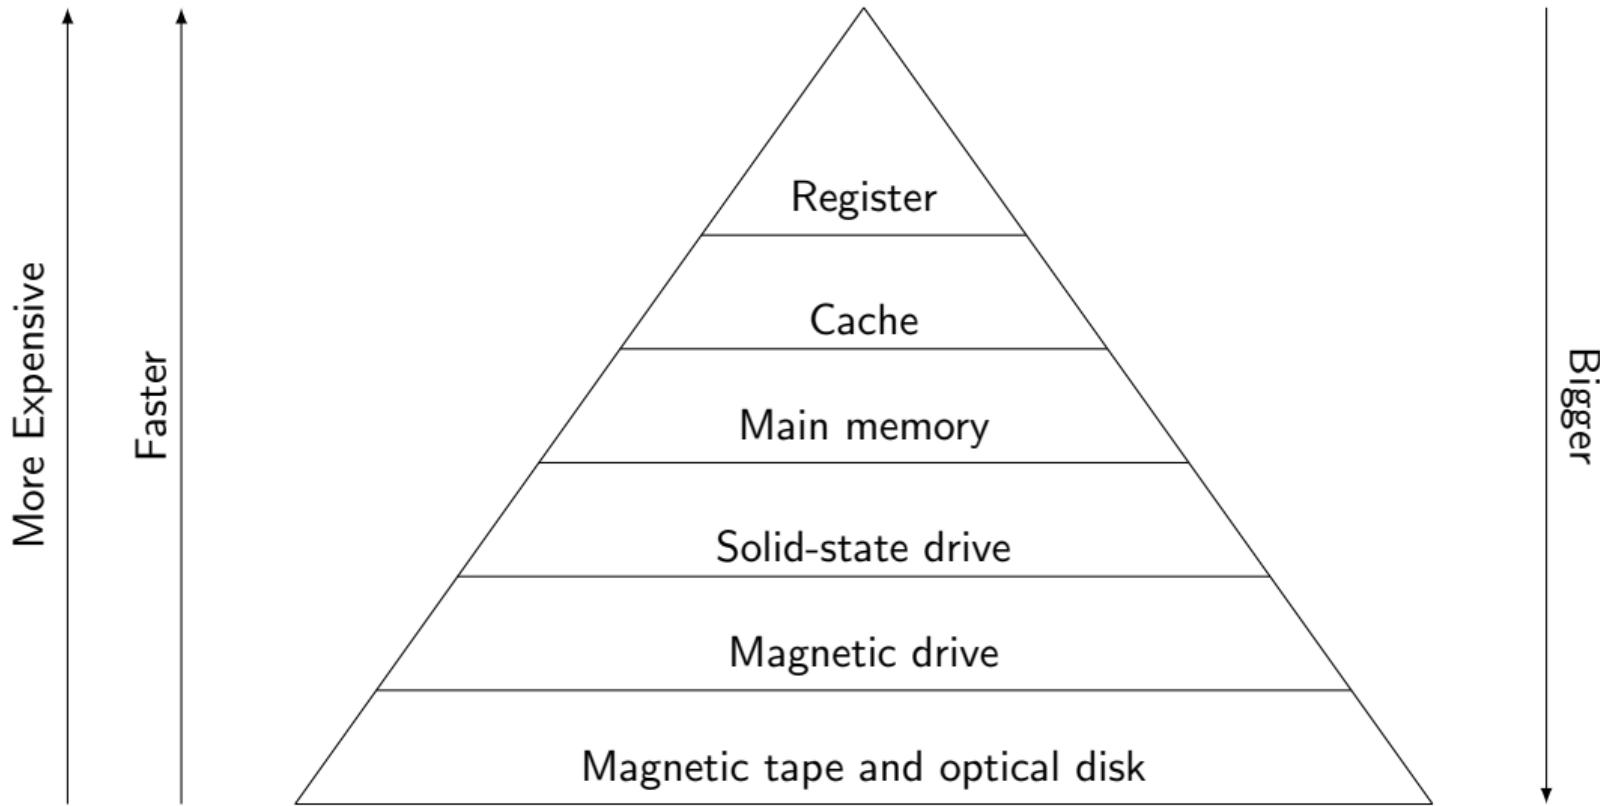


image source: wikimedia.org

# Why Multiple Storage Devices?



# Layers from Applications to the Hardware

Applications	(Web Browser, Compiler, Text editor, ...)
Operating System	(GNU/Linux, NetBSD, Windows, MacOS, ...)
Hardware	(CPU, Memory, I/O, ...)

The layer between the software and the hardware is called the Instruction Set Architecture (ISA).

# ISA (Instruction Set Architecture)

- ▶ Is an abstract model of a computer
- ▶ A.k.a. computer architecture
- ▶ Defines the set of available machine instructions (ex: MOV, POP)
- ▶ Defines the location of data (ex: registers, main memory)
- ▶ Two main categories
  - ▶ Complex Instruction Set Computer (CISC) such as Motorola 6800, Intel 8080, Intel x86 32-bit, Inter x86 64-bit, ...
  - ▶ Reduced Instruction Set Computer (RISC) such as ARM, Atmel AVR, Blackfin, i860, i960, M88000, MIPS, PA-RISC, ...
- ▶ Compatibility for a given ISA enables to reuse binaries

# ISA in SVEM

In SVEM we use Intel's x86 64-bit ISA.

Where is the ARM's ISA used quite a lot nowadays?

# x86 64-bit CPU Registers

64-bit register	Lower 32 bits	Lower 16 bits	Lower 8 bits
rax	eax	ax	al
rbx	ebx	bx	bl
rcx	ecx	cl	
rdx	edx	dx	dl
rsi	esi	si	sil
rdi	edi	di	dil
rbp	ebp	bp	bpl
rsp	esp	sp	spl
r8	r8d	r8w	r8b
r9	r9d	r9w	r9b
r10	r10d	r10w	r10b
r11	r11d	r11w	r11b
r12	r12d	r12w	r12b
r13	r13d	r13w	r13b
r14	r14d	r14w	r14b
r15	r15d	r15w	r15b

## x86 64-bit CPU Registers (cont)

- ▶ rip (instruction pointer, 32-bit called eip).
- ▶ rflags register (32-bit called flags)
- ▶ floating-point registers: eight 80-bit x87 registers: ST(0) to ST(7), eight 64-bit MMX registers: MMX0 to MMX7. (these last 8 registers overlap with the x87 registers.)
- ▶ sixteen 128-bit SSE<sup>1</sup> registers (8 for 32-bit): XMM0 to XMM15

---

<sup>1</sup>Streaming SIMD Extensions (SSE) is a single instruction, multiple data (SIMD) instruction set extension. A SIMD performs the same operation on multiple data points at the same time, thus it potentially reduces the execution time of the program.

# Adressing Modes

# Compiling & Executing a Program

- ▶ Generally a program is written in a "high-level" language (ex: C, C++, Java)
- ▶ The program is compiled to machine code
- ▶ The program is executed by the CPU
- ▶ Any operation (ex: addition) is performed using CPU registers

# Source Code, Machine Code & Assembly Code

```
int main(int argc, char** argv) {  
    char elements[100]; int sum = 0;  
    for (int i = 0; i < 100; i++) {  
        sum += elements[i];  
    }  
    return sum;  
}
```

	0000000000001125 <main>:	
1125:	55	push %rbp
1126:	48 89 e5	mov %rsp,%rbp
1129:	48 83 ec 08	sub \$0x8,%rsp
112d:	89 7d 8c	mov %edi,-0x74(%rbp)
1130:	48 89 75 80	mov %rsi,-0x80(%rbp)
1134:	c7 45 fc 00 00 00 00	movl \$0x0,-0x4(%rbp)
113b:	c7 45 f8 00 00 00 00	movl \$0x0,-0x8(%rbp)
1142:	eb 14	jmp 1158 <main+0x33>
1144:	8b 45 f8	mov -0x8(%rbp),%eax
1147:	48 98	cltq
1149:	0f b6 44 05 90	movzbl -0x70(%rbp,%rax,1),%eax
114e:	0f be c0	movsbl %al,%eax
1151:	01 45 fc	add %eax,-0x4(%rbp)
1154:	83 45 f8 01	addl \$0x1,-0x8(%rbp)
1158:	83 7d f8 63	cmpl \$0x63,-0x8(%rbp)
115c:	7e e6	jle 1144 <main+0x1f>
115e:	8b 45 fc	mov -0x4(%rbp),%eax
1161:	c9	leaveq
1162:	c3	retq
1163:	66 2e 0f 1f 84 00 00	nopw %cs:0x0(%rax,%rax,1)
116a:	00 00 00	
116d:	0f 1f 00	nopl (%rax)

- ▶ Source code (top left) and assembly code (right) are understandable for a human
- ▶ Machine code (middle) is understandable for a processor

# A Note on Assembly Syntax (Intel vs. AT&T

	Intel	AT&T
Sigils	assembler detects type of symbols	Immediate values prefixed with a "\$", registers prefixed with "%"
Parameter Order	ex: mov rsp, 42	ex: mov \$42, %rsp
Parameter Size	derived from the register name: 64-, 32-, 16-, 8-bit imply q, l, w, b, respectively ex: add ebx, 3	mnemonics are suffixed by a letter: 'q' for qword, 'l' for long (dword), 'w' for word, and 'b' for byte ex: addl \$3, %ebx
Effective Address	mov rax, [ecx + ebx*31 + memory_location] size keywords like byte, word, or dword have to be used if the size cannot be determined from the operands ex: movzb eax, [rbp + rax*1 - 0x70]	movq memory_location(%ecx,%ebx,31), %rax DISPLACEMENT(BASE,INDEX,SCALE)  ex: movzbl -0x70(%rbp,%rax,1),%eax

# Stack, RSP

```
int main(int argc, char** argv) {  
    char elements[100]; int sum = 0;  
    for (int i = 0; i < 100; i++) {  
        sum += elements[i];  
    }  
    return sum;  
}
```

- ▶ function parameters: rdi, rsi, rdx, rcx, r8, r9.  
This is the calling convention of the System V  
AMD64 Application Binary Interface (ABI).
- ▶ Variables sum, i and elements are stored on the  
stack ?
- ▶ We do not see that 100 bytes have been reserved
- ▶ Return value: eax

```
000000000000001125 <main>:  
1125    push    %rbp  
1126    mov     %rsp,%rbp  
1129    sub     $0x8,%rsp  
112d    mov     %edi,-0x74(%rbp)  
1130    mov     %rsi,-0x80(%rbp)  
1134    movl   $0x0,-0x4(%rbp)  
113b    movl   $0x0,-0x8(%rbp)  
1142    jmp    1158 <main+0x33>  
1144    mov     -0x8(%rbp),%eax  
1147    cltq  
1149    movzbl -0x70(%rbp,%rax,1),%eax  
114e    movsbl %al,%eax  
1151    add     %eax,-0x4(%rbp)  
1154    addl   $0x1,-0x8(%rbp)  
1158    cmpl   $0x63,-0x8(%rbp)  
115c    jle    1144 <main+0x1f>  
115e    mov     -0x4(%rbp),%eax  
1161    leaveq  
1162    retq  
1163    nopw   %cs:0x0(%rax,%rax,1)  
116a    nopl   (%rax)  
116d    nopl   (%rax)
```

# Stack, RSP

```
int function() { }

int main(int argc, char** argv) {
    char elements[100]; int sum = 0;
    for (int i = 0; i < 100; i++) {
        sum += elements[i];
        function();
    }
    return sum;
}
```

- ▶ main now calls function
- ▶ now, we can see that 100+ bytes have been "reserved"

```
00000000000000112c <main>:
112c    push    %rbp
112d    mov     %rsp,%rbp
1130    add    $0xfffffffffffffff80,%rsp
1134    mov     %edi,-0x74(%rbp)
1137    mov     %rsi,-0x80(%rbp)
113b    movl   $0x0,-0x4(%rbp)
1142    movl   $0x0,-0x8(%rbp)
1149    jmp    1169 <main+0x3d>
114b    mov    -0x8(%rbp),%eax
114e    cltq
1150    movzbl -0x70(%rbp,%rax,1),%eax
1155    movsbl %al,%eax
1158    add    %eax,-0x4(%rbp)
115b    mov    $0x0,%eax
1160    callq  1125 <function>
1165    addl   $0x1,-0x8(%rbp)
1169    cmpl   $0x63,-0x8(%rbp)
116d    jle    114b <main+0x1f>
116f    mov    -0x4(%rbp),%eax
1172    leaveq
1173    retq
1174    nopw   %cs:0x0(%rax,%rax,1)
```

# Optimization 1

```
int main(int argc, char** argv) {  
    char elements[100]; int sum = 0;  
    for (int i = 0; i < 100; i++) {  
        sum += elements[i];  
    }  
    return sum;  
}
```

- ▶ Uses registers instead of the stack
- ▶ gcc -O1

00000000000001126	<main>:
1126	lea -0x70(%rsp),%rdx
112b	lea 0x64(%rdx),%rsi
112f	mov \$0x0,%eax
1134	movsb (%rdx),%ecx
1137	add %ecx,%eax
1139	add \$0x1,%rdx
113d	cmp %rsi,%rdx
1140	jne 1134 <main+0xe>
1142	retq
1143	nopw %cs:0x0(%rax,%rax,1)
114a	
114d	nopl (%rax)

# Optimization 2

```
int main(int argc, char** argv) {  
    char elements[100]; int sum = 0;  
    for (int i = 0; i < 100; i++) {  
        sum += elements[i];  
    }  
    return sum;  
}
```

- ▶ Improve speed by changing some instructions
- ▶ gcc -O2

	00000000000000001040	<main>:
1040	lea	-0x78(%rsp),%rdx
1045	xor	%eax,%eax
1047	lea	0x64(%rdx),%rsi
104b	nopl	0x0(%rax,%rax,1)
1050	movsbl	(%rdx),%ecx
1053	add	\$0x1,%rdx
1057	add	%ecx,%eax
1059	cmp	%rdx,%rsi
105c	jne	1050 <main+0x10>
105e	retq	
105f	nop	

# Optimization 1 vs. Optimization 2

```
00000000000000001126 <main>:  
1126    lea    -0x70(%rsp),%rdx  
112b    lea    0x64(%rdx),%rsi  
112f    mov    $0x0,%eax  
1134    movsb  (%rdx),%ecx  
1137    add    %ecx,%eax  
1139    add    $0x1,%rdx  
113d    cmp    %rsi,%rdx  
1140    jne    1134 <main+0xe>  
1142    retq  
1143    nopw  %cs:0x0(%rax,%rax,1)  
114a    nopl  (%rax)  
114d    nopl  (%rax)
```

```
00000000000000001040 <main>:  
1040    lea    -0x78(%rsp),%rdx  
1045    xor    %eax,%eax  
1047    lea    0x64(%rdx),%rsi  
104b    nopl  0x0(%rax,%rax,1)  
1050    movsb  (%rdx),%ecx  
1053    add    $0x1,%rdx  
1057    add    %ecx,%eax  
1059    cmp    %rdx,%rsi  
105c    jne    1050 <main+0x10>  
105e    retq  
105f    nop
```

# Optimization 3

```
int main(int argc, char** argv) {  
    char elements[100]; int sum = 0;  
    for (int i = 0; i < 100; i++) {  
        sum += elements[i];  
    }  
    return sum;  
}
```

- ▶ Speed optimization
- ▶ gcc -O3
- ▶ (full assembly code on the next slide)
- ▶ What is the "idea" behind this optimization?

```
000000000000001040 <main>:  
1040 pxor %xmm2,%xmm2  
1044 movdqa -0x78(%rsp),%xmm3  
104a movsbl -0x18(%rsp),%edx  
104f movdqa %xmm2,%xmm1  
1053 pcmpgtb %xmm3,%xmm1  
1057 movdqa %xmm3,%xmm0  
105b punpcklbw %xmm1,%xmm0  
105f punpckhbw %xmm1,%xmm3  
1063 pxor %xmm1,%xmm1  
1067 movdqa %xmm1,%xmm4  
106b movdqa %xmm0,%xmm5  
106f pcmpgtw %xmm0,%xmm4  
1073 punpcklwd %xmm4,%xmm5  
1077 punpckhwd %xmm4,%xmm0  
107b movdqa %xmm1,%xmm4  
107f pcmpgtw %xmm3,%xmm4  
1083 paddd %xmm3,%xmm0  
1087 movdqa %xmm3,%xmm5  
108b punpcklwd %xmm4,%xmm5  
108f punpckhwd %xmm4,%xmm3  
[...]
```

# Optimization 3 (cont)

00000000000000001040	<code>main:</code>	10cd	paddd %xmm6,%xmm0	115d	pcmpgtw %xmm4,%xmm5	11ed	movdqa %xmm3,%xmm4
1040	pxor %xmm2,%xmm2	10d1	paddd %xmm4,%xmm0	1161	movdqa %xmm4,%xmm6	11f1	punpckhbw %xmm2,%xmm3
1044	movdqa -0x78(%rsp),%xmm3	10d5	movdqa %xmm1,%xmm4	1165	punpcklwd %xmm5,%xmm6	11f5	punpcklbw %xmm2,%xmm4
104a	movsbl -0x18(%rsp),%edx	10d9	pcmpgtw %xmm3,%xmm4	1169	punpckhwd %xmm5,%xmm4	11f9	movdqa %xmm3,%xmm2
104f	movdqa %xmm2,%xmm1	10dd	punpcklwd %xmm4,%xmm5	116d	movdqa %xmm3,%xmm5	11fd	movdqa %xmm1,%xmm3
1053	pcmpgtb %xmm3,%xmm1	10e1	punpckhwd %xmm4,%xmm3	1171	paddd %xmm6,%xmm0	1201	movdqa %xmm4,%xmm5
1057	movdqa %xmm3,%xmm0	10e5	paddd %xmm5,%xmm0	1175	paddd %xmm4,%xmm0	1205	pcmpgtw %xmm4,%xmm3
105b	punpcklbw %xmm1,%xmm0	10e9	movdqa %xmm2,%xmm5	1179	movdqa %xmm1,%xmm4	1209	pcmpgtw %xmm2,%xmm1
105f	punpckhbw %xmm1,%xmm3	10ed	paddd %xmm3,%xmm0	117d	pcmpgtw %xmm3,%xmm4	120d	punpcklwd %xmm3,%xmm5
1063	pxor %xmm1,%xmm1	10f1	movdqa -0x58(%rsp),%xmm3	1181	punpcklwd %xmm4,%xmm5	1211	punpckhw %xmm3,%xmm4
1067	movdqa %xmm1,%xmm4	10f7	pcmpgtb %xmm3,%xmm5	1185	punpckhwd %xmm4,%xmm3	1215	movdqa %xmm2,%xmm3
106b	movdqa %xmm0,%xmm5	10fb	movdqa %xmm3,%xmm4	1189	paddd %xmm5,%xmm0	1219	paddd %xmm5,%xmm0
106f	pcmpgtw %xmm0,%xmm4	10ff	punpcklbw %xmm5,%xmm4	118d	movdqa %xmm2,%xmm5	121d	punpcklwd %xmm1,%xmm3
1073	punpcklwd %xmm4,%xmm5	1103	punpckhbw %xmm5,%xmm3	1191	paddd %xmm3,%xmm0	1221	punpckhw %xmm1,%xmm2
1077	punpckhwd %xmm4,%xmm0	1107	movdqa %xmm1,%xmm5	1195	movdqa -0x38(%rsp),%xmm3	1225	paddd %xmm4,%xmm0
107b	movdqa %xmm1,%xmm4	110b	pcmpgtw %xmm4,%xmm5	119b	pcmpgtb %xmm3,%xmm5	1229	paddd %xmm3,%xmm0
107f	pcmpgtw %xmm3,%xmm4	110f	movdqa %xmm4,%xmm6	119f	movdqa %xmm3,%xmm4	122d	paddd %xmm2,%xmm0
1083	paddd %xmm5,%xmm0	1113	punpcklwd %xmm5,%xmm6	11a3	punpcklbw %xmm5,%xmm4	1231	movdqa %xmm0,%xmm1
1087	movdqa %xmm3,%xmm5	1117	punpckhwd %xmm5,%xmm4	11a7	punpckhbw %xmm5,%xmm3	1235	psrldq \$0x8,%xmm1
108b	punpcklwd %xmm4,%xmm5	111b	movdqa %xmm3,%xmm5	11ab	movdqa %xmm1,%xmm5	123a	paddd %xmm1,%xmm0
108f	punpckhwd %xmm4,%xmm3	111f	paddd %xmm6,%xmm0	11af	pcmpgtw %xmm4,%xmm5	123e	movdqa %xmm0,%xmm1
1093	paddd %xmm5,%xmm0	1123	paddd %xmm4,%xmm0	11b3	movdqa %xmm4,%xmm6	1242	psrldq \$0x4,%xmm1
1097	movdqa %xmm2,%xmm5	1127	movdqa %xmm1,%xmm4	11b7	punpcklwd %xmm5,%xmm6	1247	paddd %xmm1,%xmm0
109b	paddd %xmm3,%xmm0	112b	pcmpgtw %xmm3,%xmm4	11bb	punpckhwd %xmm5,%xmm4	124b	movd %xmm0,%eax
109f	movdqa -0x68(%rsp),%xmm3	112f	punpcklwd %xmm4,%xmm5	11bf	movdqa %xmm3,%xmm5	124f	add %edx,%eax
10a5	pcmpgtb %xmm3,%xmm5	1133	punpckhwd %xmm4,%xmm3	11c3	paddd %xmm6,%xmm0	1251	movsbl -0x17(%rsp),%edx
10a9	movdqa %xmm3,%xmm4	1137	paddd %xmm5,%xmm0	11c7	paddd %xmm4,%xmm0	1256	add %eax,%edx
10ad	punpcklbw %xmm5,%xmm4	113b	movdqa %xmm2,%xmm5	11cb	movdqa %xmm1,%xmm4	1258	movsbl -0x16(%rsp),%eax
10b1	punpckhbw %xmm5,%xmm3	113f	paddd %xmm3,%xmm0	11cf	pcmpgtw %xmm3,%xmm4	125d	add %eax,%edx
10b5	movdqa %xmm1,%xmm5	1143	movdqa -0x48(%rsp),%xmm3	11d3	punpcklwd %xmm4,%xmm5	125f	movsbl -0x15(%rsp),%eax
10b9	pcmpgtw %xmm4,%xmm5	1149	pcmpgtb %xmm3,%xmm5	11d7	punpckhwd %xmm4,%xmm3	1264	add %edx,%eax
10bd	movdqa %xmm4,%xmm6	114d	movdqa %xmm3,%xmm4	11db	paddd %xmm5,%xmm0	1266	retq
10c1	punpcklwd %xmm5,%xmm6	1151	punpcklbw %xmm5,%xmm4	11df	paddd %xmm3,%xmm0	1267	nopw 0x0(%rax,%rax,1)
10c5	punpckhwd %xmm5,%xmm4	1155	punpckhbw %xmm5,%xmm3	11e3	movdqa -0x28(%rsp),%xmm3	126e	
10c9	movdqa %xmm3,%xmm5	1159	movdqa %xmm1,%xmm5	11e9	pcmpgtb %xmm3,%xmm2		

# Function Call and the Stack

```
int function3() {
    char buffer[0x42];
    return 3;
}
int function2() {
    char buffer[0x42];
    function3();
    return 3;
}
int function1() {
    int a = 2;
    function2();
}
int main(int argc, char** argv) {
    function1();
}
```

- ▶ rsp and rbp registers to keep the context
- ▶ rsp, Stack Pointer
- ▶ rbp, Base Pointer, base address for accessing local variables on the stack

# Function Call and the Stack

```
00000000000000001125 <function3>:  
1125    push    %rbp  
1126    mov     %rsp,%rbp  
1129    mov     $0x3,%eax  
112e    pop     %rbp  
112f    retq
```

```
00000000000000001130 <function2>:  
1130    push    %rbp  
1131    mov     %rsp,%rbp  
1134    sub     $0x50,%rsp  
1138    mov     $0x0,%eax  
113d    callq   1125 <function3>  
1142    mov     $0x3,%eax  
1147    leaveq  
1148    retq
```

```
00000000000000001149 <function1>:  
1149    push    %rbp  
114a    mov     %rsp,%rbp  
114d    sub     $0x10,%rsp  
1151    movl   $0x2,-0x4(%rbp)  
1158    mov     $0x0,%eax
```

```
115d    callq   1130 <function2>  
1162    nop  
1163    leaveq  
1164    retq
```

```
00000000000000001165 <main>:  
1165    push    %rbp  
1166    mov     %rsp,%rbp  
1169    sub     $0x10,%rsp  
116d    mov     %edi,-0x4(%rbp)  
1170    mov     %rsi,-0x10(%rbp)  
1174    mov     $0x0,%eax  
1179    callq   1149 <function1>  
117e    mov     $0x0,%eax  
1183    leaveq  
1184    retq  
1185    nopw   %cs:0x0(%rax,%rax,1)  
118c  
118f    nop
```

Note: leave  $\Leftrightarrow$  mov %rbp, %rsp; pop %rbp;  
Note: enter N, vs. leave

# Byte Alignment

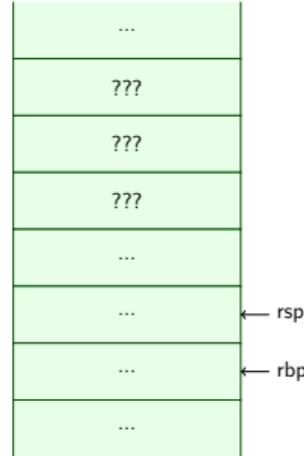
- ▶ The x86-64 System V ABI requires 16 bytes stack alignment before a call.
- ▶ For main, the compiler puts the two arguments (argc and argv) on the stack. The first argument is an int (4 bytes) the second a 64-bit pointer (8 bytes). The required size for these arguments is thus 12 bytes. Because of the 16 byte stack alignment, the compiler allocates 16 bytes (0x10)
- ▶ For function1, the reasoning is the same with a single int local variable (4 bytes). The compiler allocates 16 bytes (0x10) on the stack instead of 4 to respect the 16 bytes stack alignment.
- ▶ For function2, the compiler allocates 80 bytes on the stack (0x50) instead of 66 bytes (0x42) to respect the 16 bytes stack alignment.
- ▶ Note that the return address and rbp pushed on the stack form a 16 byte block (already 16 byte aligned)

## Why Such Alignment?

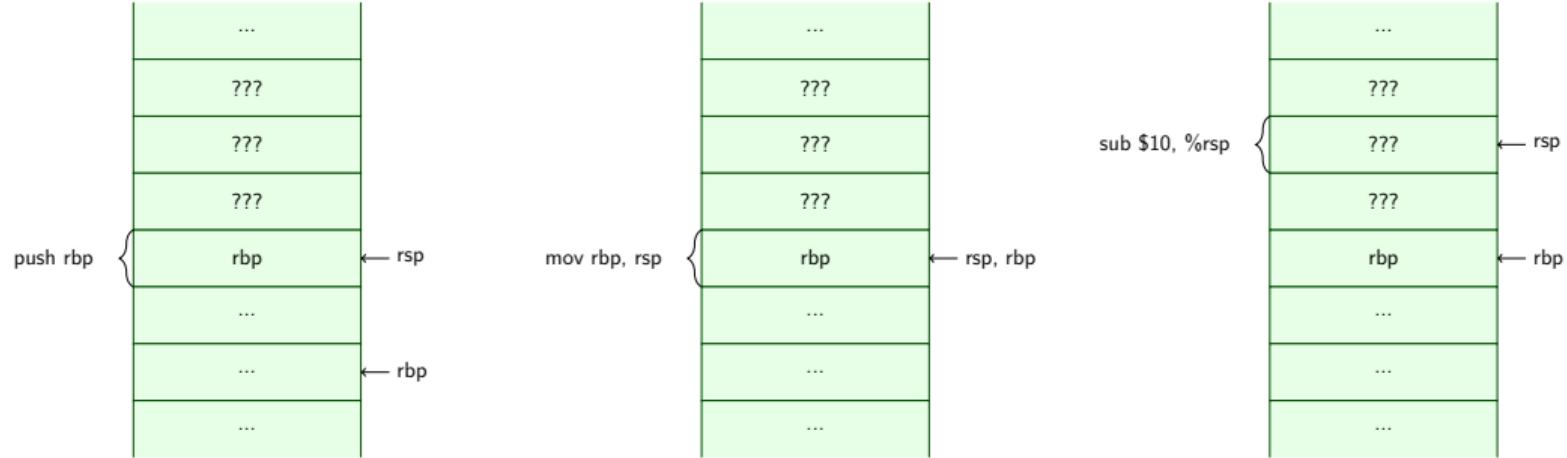
According to the System V AMD64 ABI specifications, the alignment is – at least – required for some instructions:

*"The alignment requirement allows the use of SSE instructions when operating on the array. The compiler cannot in general calculate the size of a variable-length array (VLA), but it is expected that most VLAs will require at least 16 bytes, so it is logical to mandate that VLAs have at least a 16-byte alignment."*

In main...



# In main... (cont)



# Exercise

Draw the stack at the following points

- ▶ Just before function1 is called
- ▶ Just before function2 is called
- ▶ Just before function3 is called

# Von Neumann Architecture vs. Harvard Architecture

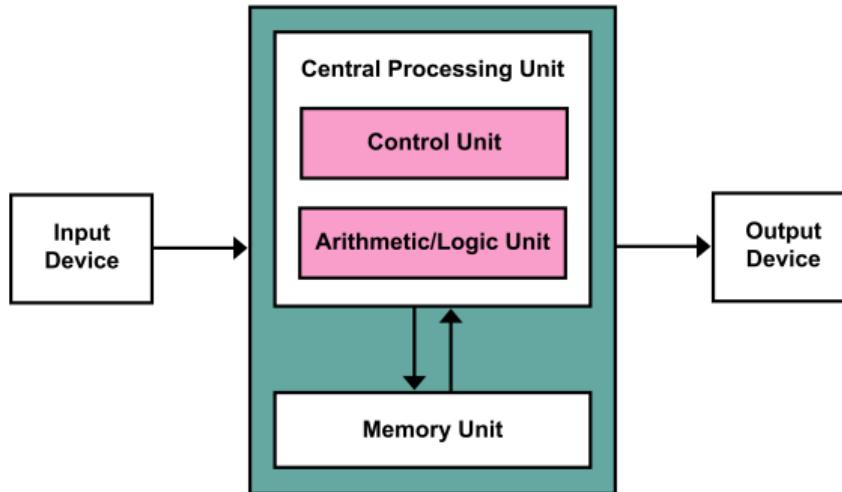


image source: wikimedia.org

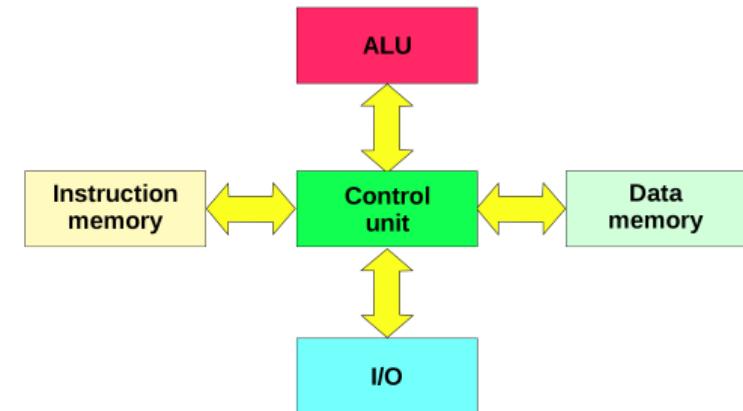


image source: wikimedia.org

# Sequential Execution Model

- 0 Fetch
- 1 Decode
- 2 Read Inputs
- 3 Execute
- 4 Write Outputs
- 5 Next Instruction (goto 0)

## 0 Fetch, 1 Decode

- ▶ Length: variable for x86 (max 15 bytes)
- ▶ Encoding/Decoding: complex
- ▶ Hard to know where next instruction is

## 2 Read Inputs

- ▶ Where does an instruction read data from?
- ▶ A register
  - ▶ short term memory
  - ▶ faster than main memory
  - ▶ directly named in instructions
- ▶ Main memory
  - ▶ longer term memory
  - ▶ slower than registers
  - ▶ location computer from registers
- ▶ Immediate
  - ▶ values encoded within instructions

### 3 Execute

- ▶ Manipulates main memory and registers to compute a result

## 4 Write Outputs

- ▶ Where does an instruction read data from?
- ▶ A register
  - ▶ short term memory
  - ▶ faster than main memory
  - ▶ directly named in instructions
- ▶ Main memory
  - ▶ longer term memory
  - ▶ slower than registers
  - ▶ location computer from registers

- ▶ Default:  $\text{@i} + \text{sizeof(i)}$
- ▶ Branching (pc relative, absolute, indirect through register)
- ▶ Condition

Question?

# Difference of Stack Addresses in Gdb and Outside Gdb

1. Gdb adds environment variables. Since environment variables are stored on the stack, there is a difference in the addresses of local variables inside gdb and outside gdb (because the environment variables gdb adds are not present outside gdb, and thus the stack is shorter outside gdb).
2. Gdb always gives as first argument to the main function, the absolute path to the binary (ex: /home/user/lab02/test). If, outside gdb, you start the program with ./test (because you are in the directory /home/user/lab02/ containing the binary) you use relative path and thus end up with a shorter string representing the path to the binary. Since this string is stored on the stack it could have an influence of the stack size as well.

## What are environment variables?

- ▶ They are dynamically defined variables that can influence the running of programs.
- ▶ For instance, environment variable PATH, defines in which directories to search for a binary.

# Difference of Stack Addresses – Same Environment Variables (cont)

Print the Value of an Environment Variable

```
$ echo $PATH  
/home/user/bin:/usr/local/bin:/usr/bin:/bin
```

Print all Environment Variables

```
$ env  
SHELL=/bin/bash  
HOME=/home/user  
LANG=en_US.UTF-8  
TERM=screen  
USER=user  
PATH=/home/user/bin:/usr/local/bin:/usr/bin:/bin  
[...]
```

# Difference of Stack Addresses – Same Environment Variables (cont)

Remove all Environment Variables

```
$ env -i
```

Start Program test with no Environment Variable

```
$ env -i ./test
```

Start gdb with no Environment Variable (gdb adds two env. variables)

```
$ env -i gdb ./test
```

```
(gdb) show env
```

```
LINES=61
```

```
COLUMNS=92
```

Start Program test with same Environment as gdb

```
env -i LINES=XX COLUMNS=YY ./test
```

# Difference of Stack Addresses – Same Environment Variables (cont)

## Print Environment Variables in C

You can add the following code to print environment variables in your C program. That way, you can easily make sure that the environments are the same within and outside gdb.

```
#include <stdio.h>

extern char** environ;

int main(int argc, char** argv) {
    for (char **env = environ; *env != 0; env++) {
        char *env_var = *env;
        printf("' %s '\n", env_var);
    }
    return 0;
}
```

# Difference of Stack Addresses – First Parameter to main

## First Parameter to main

The first parameter to the main function is always the path to the executed binary. You can launch the binary using a relative path (ex: ./test) or an absolute path (ex: /home/user/lab02/test). Gdb always uses the absolute path. To use the absolute path (suppose you are in directory /home/user/lab02) you can execute one of the following commands:

```
$ /home/user/lab02/test
```

```
$ `pwd`/test
```

# Difference of Stack Addresses – Conclusion

## Same Stack Size in gdb and Outside gdb

To launch the program test use:

```
$ env -i LINES=XX COLUMNS=YY /home/user/lab02/test
```

To launch gdb use:

```
$ env -i gdb ./test
```

## The number of bytes is important, not the content

Note that what is important is the number of bytes of LINES=XX and COLUMNS=YY. If in gdb you have COLUMNS=123, then you should replace COLUMNS=YY by COLUMNS=YYY, otherwise the stacks will not be the same.