

# **Lecture 3**

# **Data Execution Prevention**

MICS - 2019

**Dr. Alexandre Bartel**  
[www.abartel.net](http://www.abartel.net)

# Previously...

# Lecture 1

- Software Development
- Software Security
- Software Vulnerability

# Lecture 1: Software Development

- Tukey, 1958
- Functions, compilers, documentation (vs. hardware)
- Life-cycle
  - idea, requirements, design, implementation, deployment
- Approaches: Waterfall, Agile
- Goals: less risk, better quality

# Lecture 1: Software Security

- Security policy
- Software system tries to maintain the following attributes in accordance with the security policy:
  - C...
  - I...
  - A...
- How? With security mechanisms
  - Access control
  - Sandbox

# Lecture 1: Software Vulnerability

- Life cycle:
  - Birth, discovery, disclosure, correction, publicity, scripting, death
- Non-disclosure, full disclosure, responsible disclosure
- CVE number, MITRE

# Lecture 2: Buffer Overflow on the Stack

- Buffer: consecutive bytes in memory
- Local buffer: stored on the stack
- Function f1 calls f2 at instruction i: return @ of instruction i + 1 on the stack
- Buffer overflow overwrites return @
- Attacker puts shellcode in buffer jumps to it

# “Introduction to Software Security”\_Course Plan

## 2. Memory Attacks and Defenses

- ➔ Buffer overflow
- ➔ Heap overflow
- ➔ Integer overflow
- ➔ String format vulnerabilities
- ➔ Type confusion
- ➔ Use After Free



# **Preventing Buffer Overflow Attacks On the Stack**

# Attack Prevention: First Idea

- Gently ask developers to check bounds!
- Does not work:
  - Ex: Intel [1], AMD [2]
- Problem 1: programmers might do it... or NOT
- Problem 2: does not protect legacy software

[1] Ermolov, Mark, and Maxim Goryachy. "How to Hack a Turned-Off Computer, or Running Unsigned Code in Intel Management Engine." Black Hat Europe (2017).

[2] <https://www.bleepingcomputer.com/news/security/security-flaw-in-amds-secure-chip-on-chip-processor-disclosed-online/> (6 January 2018)

# Attack Prevention: Second Idea

- Mark the stack as NON-executable
  - Called Data Execution Prevention (DEP)
  - AKA Non-eXecute bit (NX bit)
- The attacker can still put the shellcode in the buffer
- But jumping to it triggers a segmentation fault

**Problem Solved?**

# Problem Solved?

- Attacker can still “jump” anywhere he/she likes
- Attacker could put the shellcode at some executable page in memory (ex: JIT) and jump to it.
  - Problem: where is the address of the page?

# Problem Solved?

- Attacker can still “jump” anywhere he/she likes
- Attacker could execute small code snippets ending in “ret”
- Introducing “gadgets”
- Introducing “ROP” (Return Oriented Programming)

# ROP Gadgets

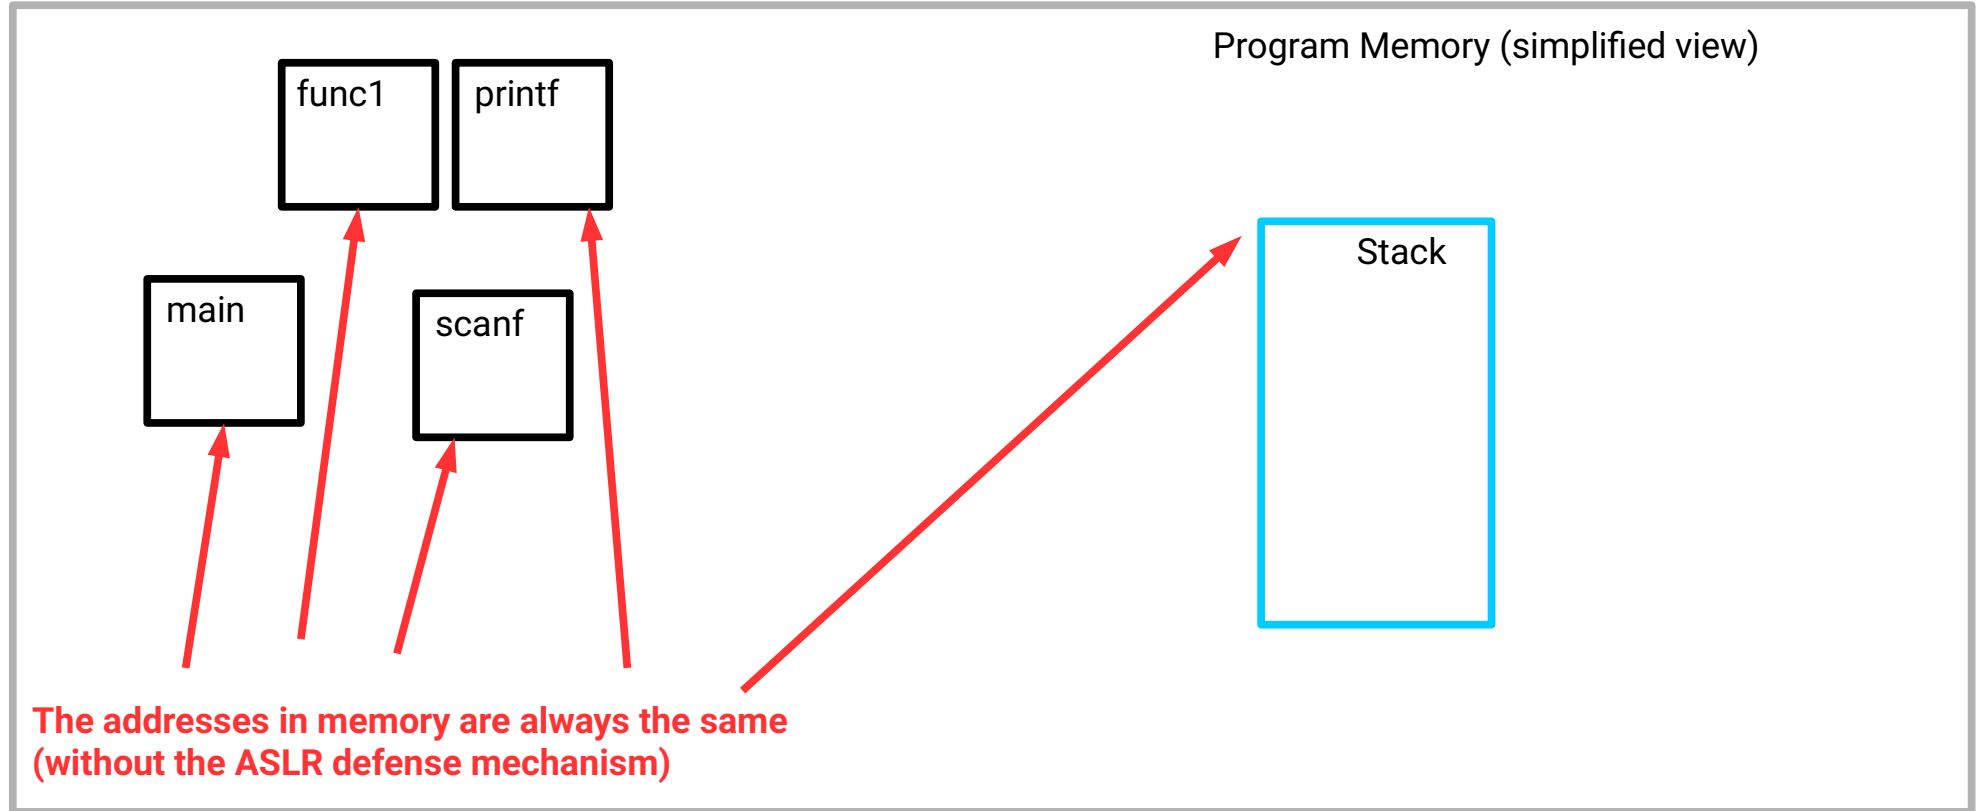
- Attacker wants to execute program P
- P features N instructions
- Recipe (simplified):
  - For instruction “i” in program “P”:
    - Find ROP gadget “Gi” executing “i”
  - Chain all ROP gadgets together with data (this as called a ROP chain)
  - Execute program “P” through ROP gadgets

# Why ROP works

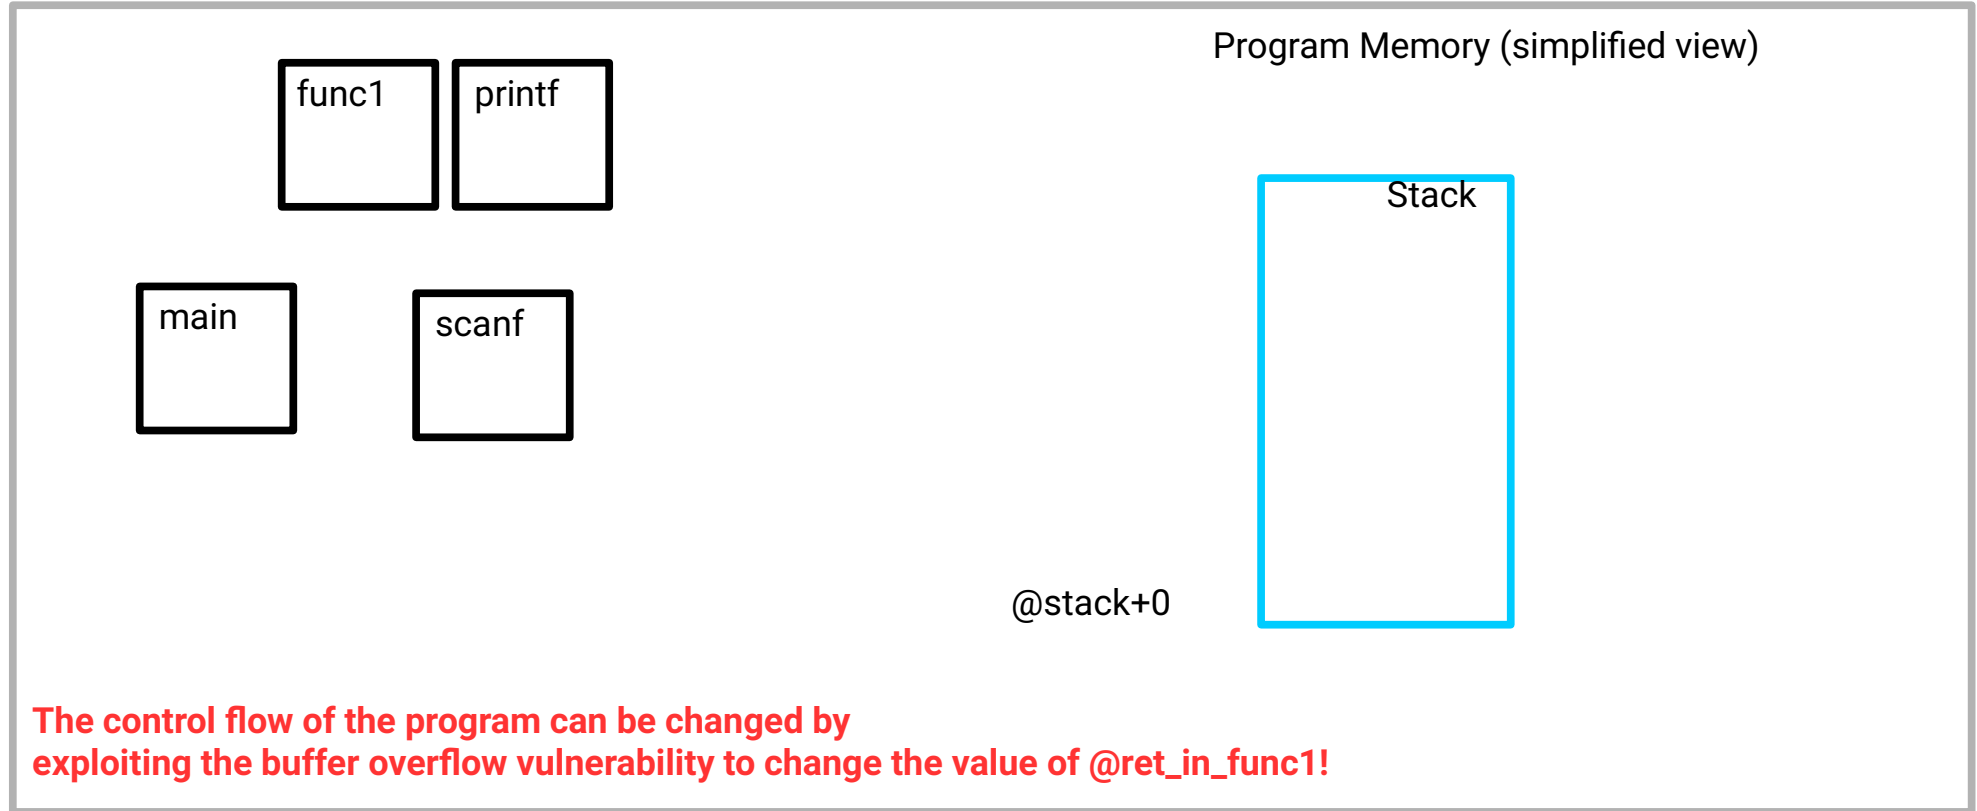
- Code is always loaded at the same address
- Only data is pushed to the stack
- Code is “reused”
- No code is injected



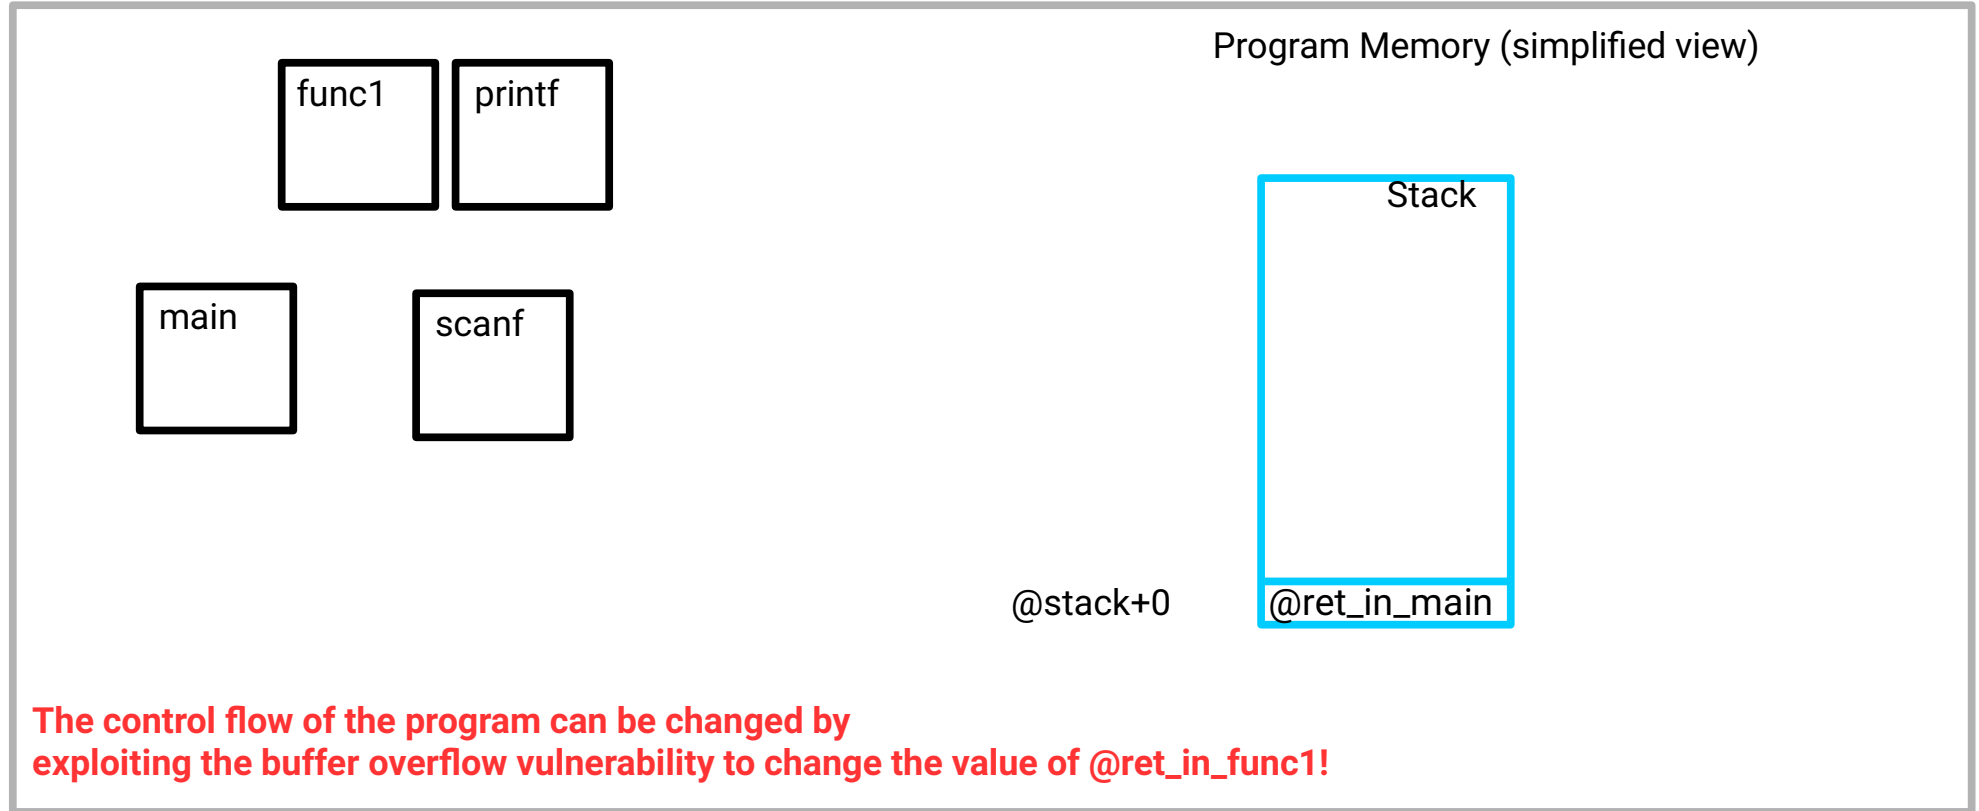
# Custom Code Execution



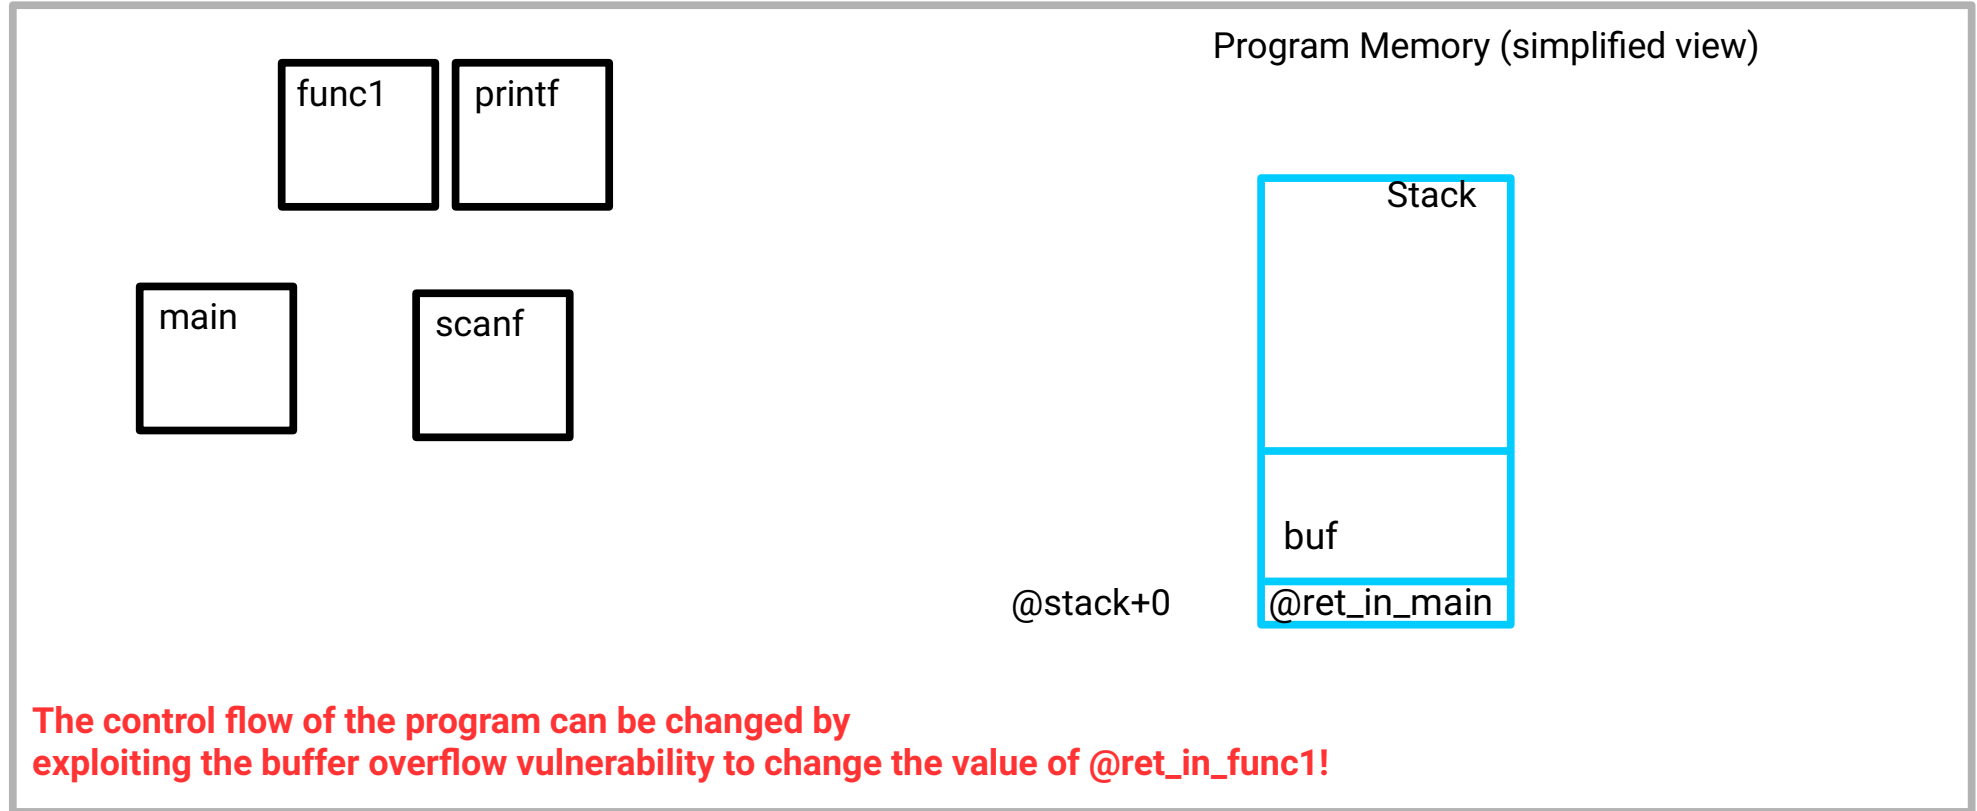
# Custom Code Execution



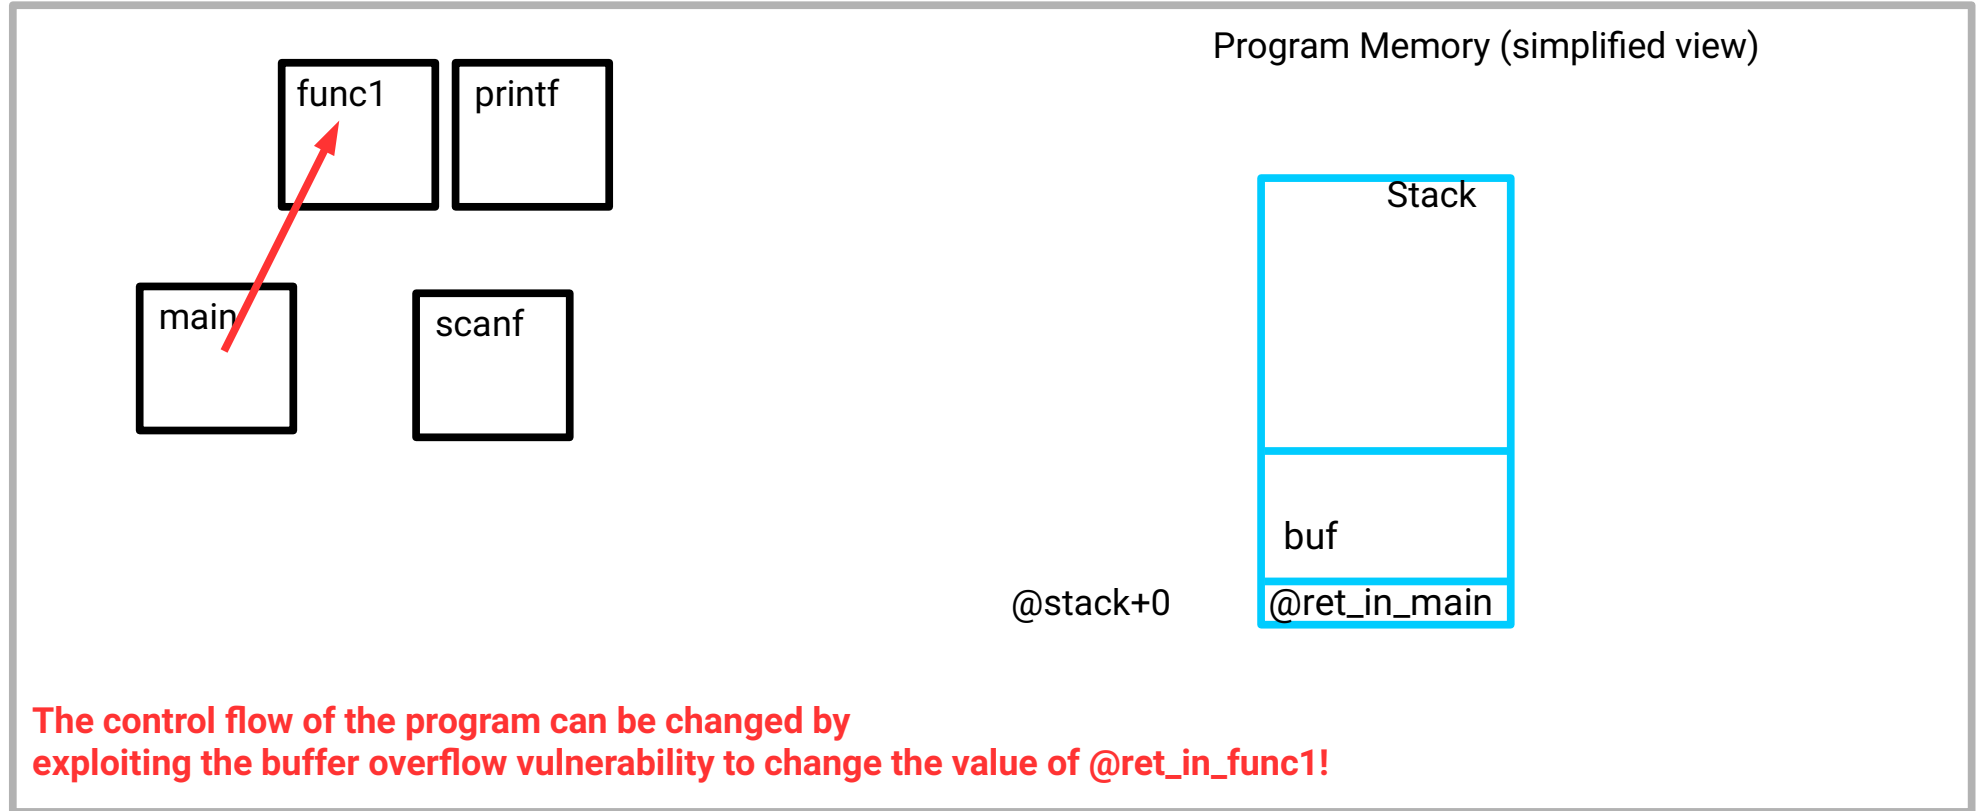
# Custom Code Execution



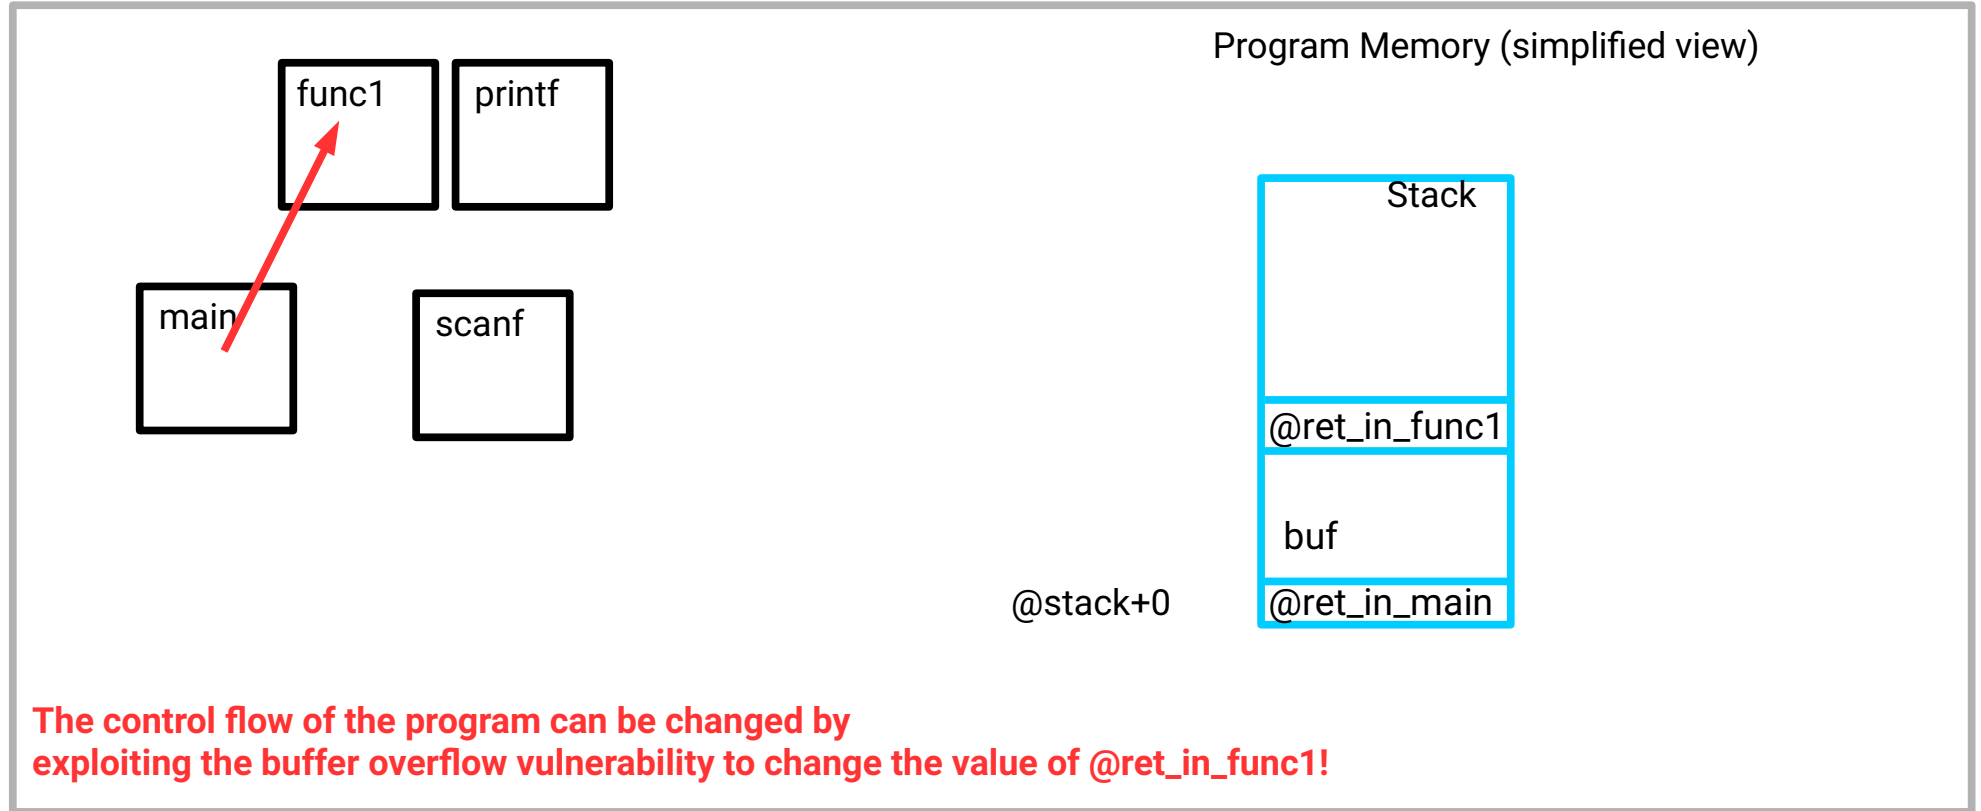
# Custom Code Execution



# Custom Code Execution



# Custom Code Execution

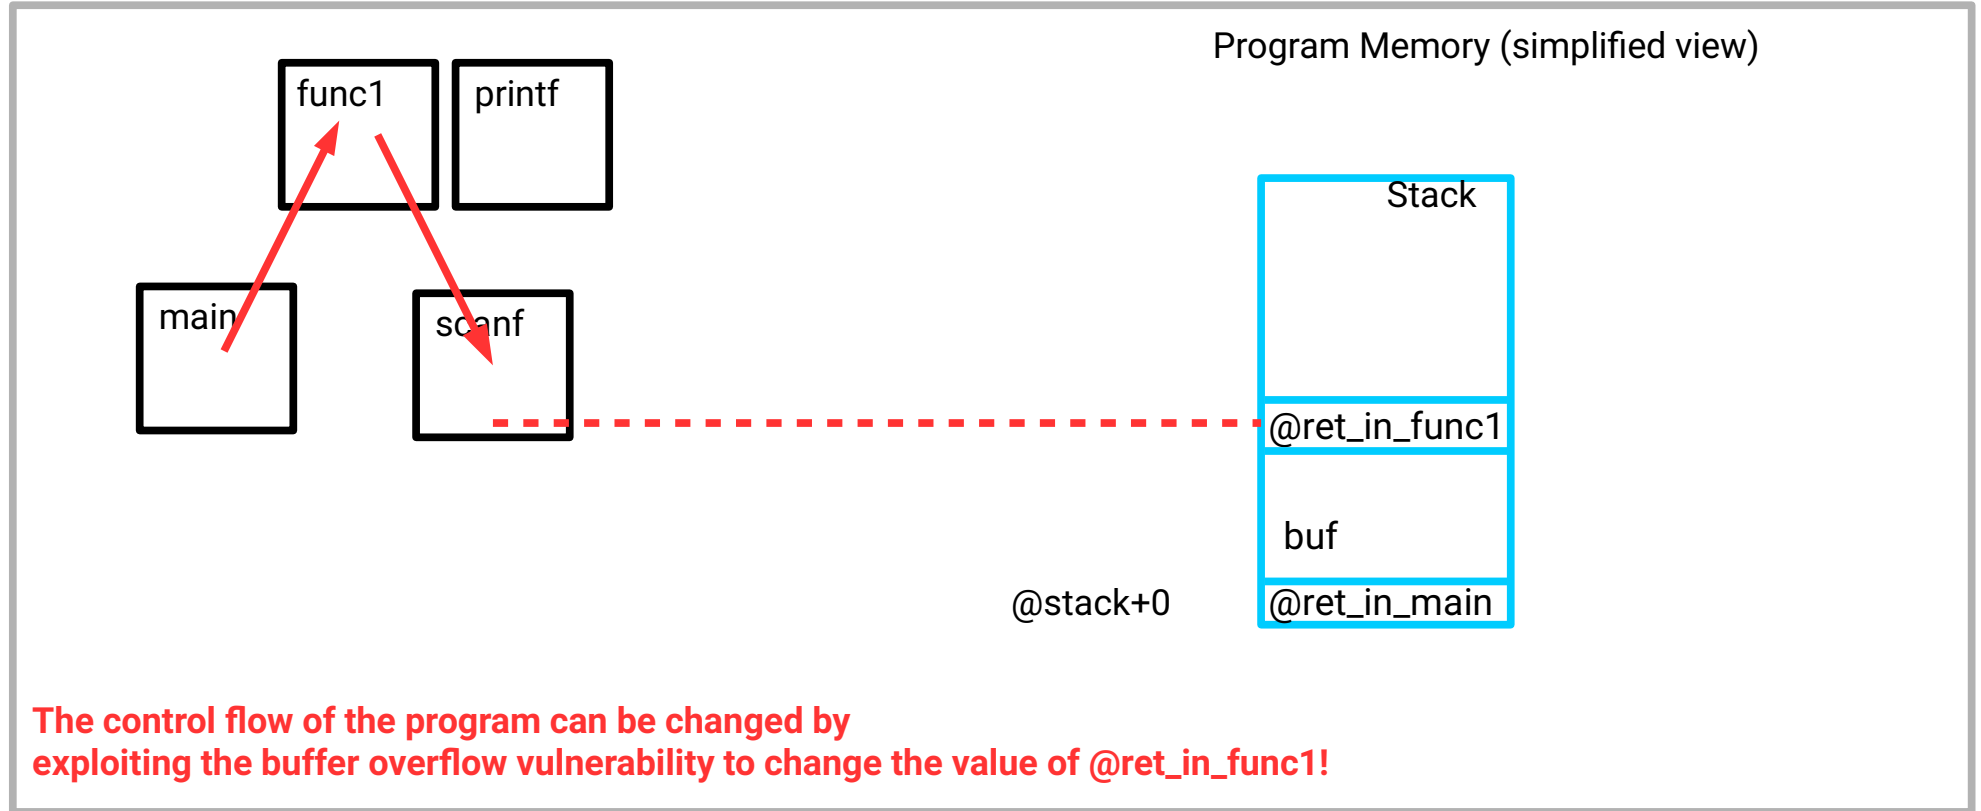


# Custom Code Execution



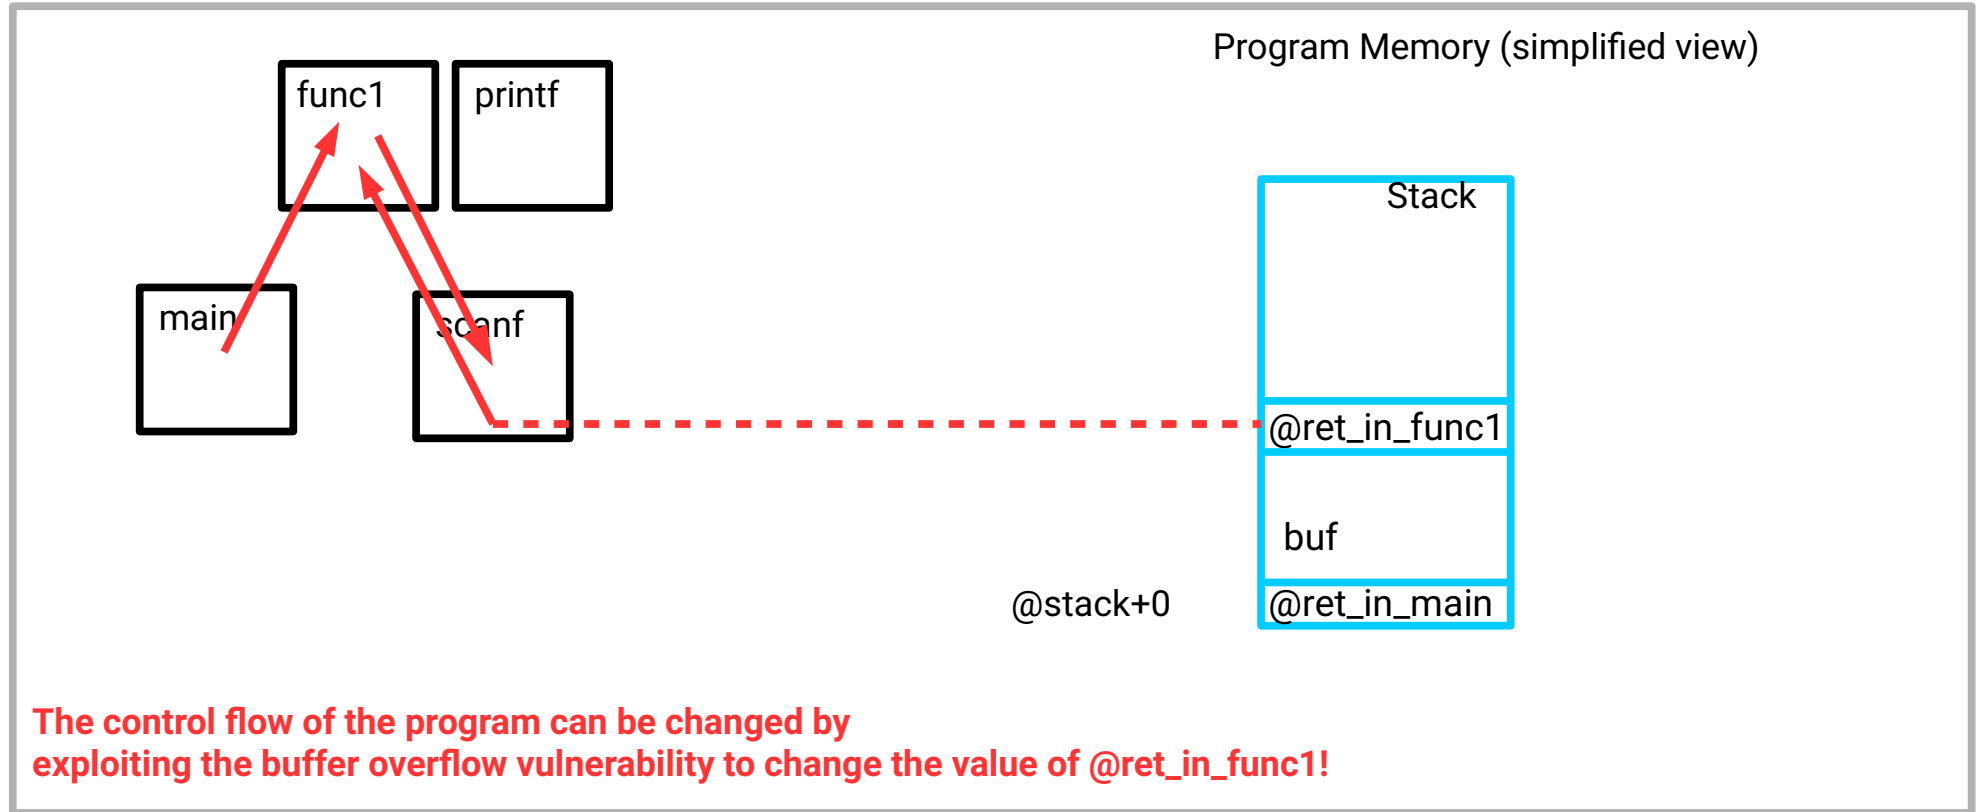
The control flow of the program can be changed by exploiting the buffer overflow vulnerability to change the value of @ret\_in\_func1!

# Custom Code Execution

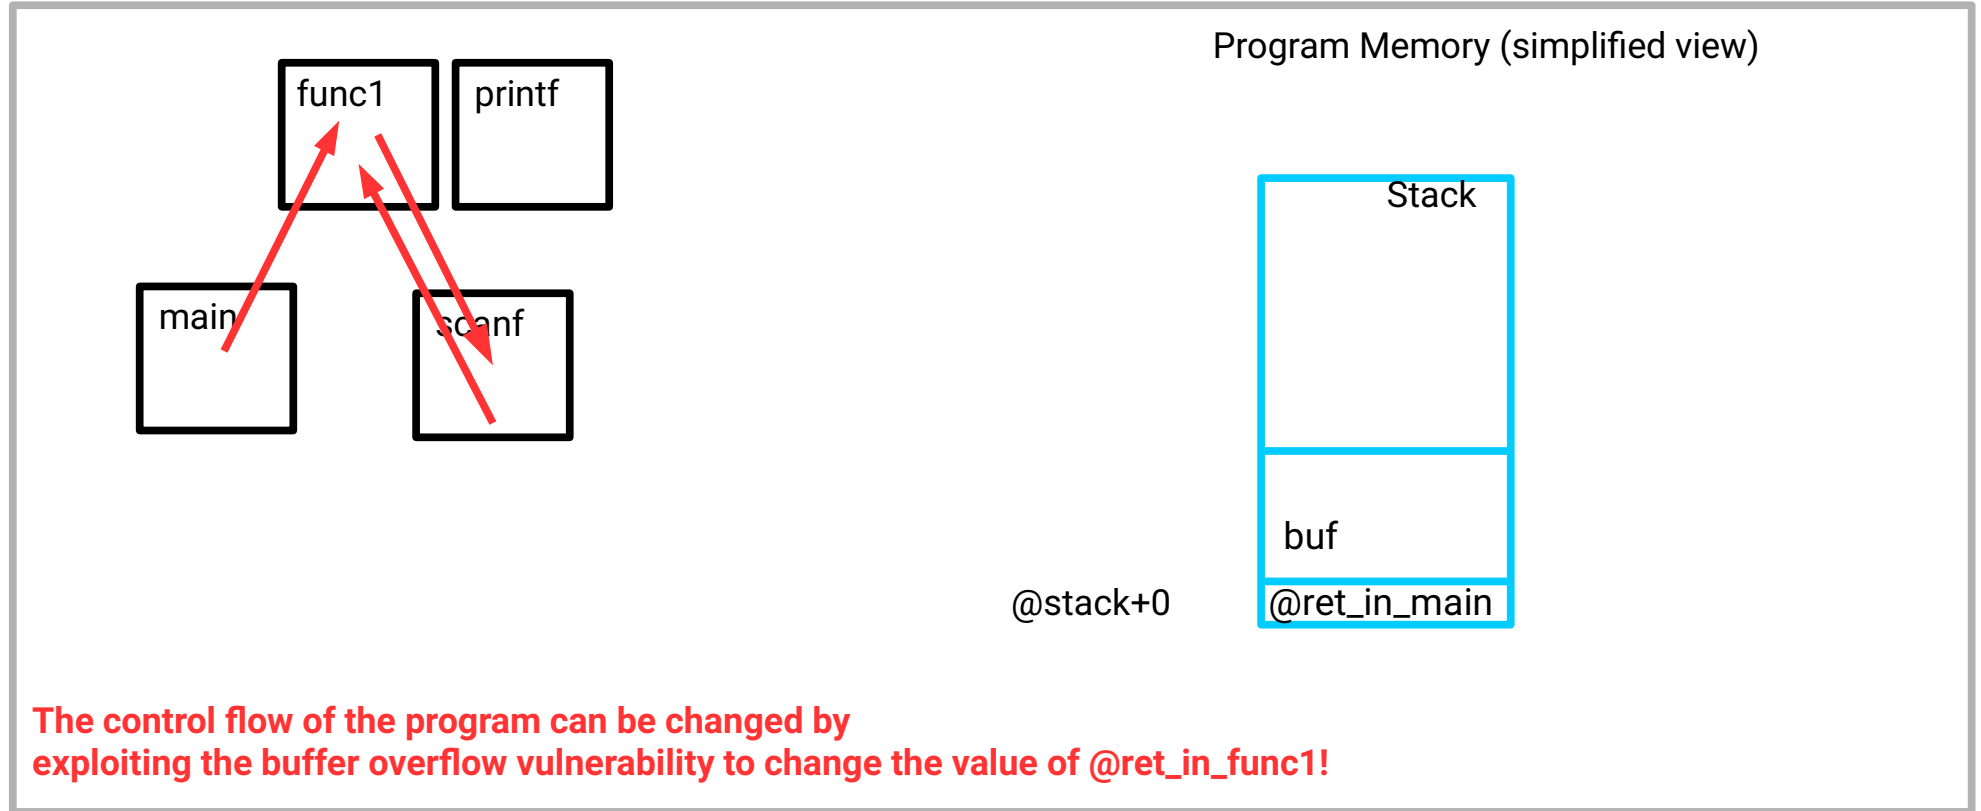




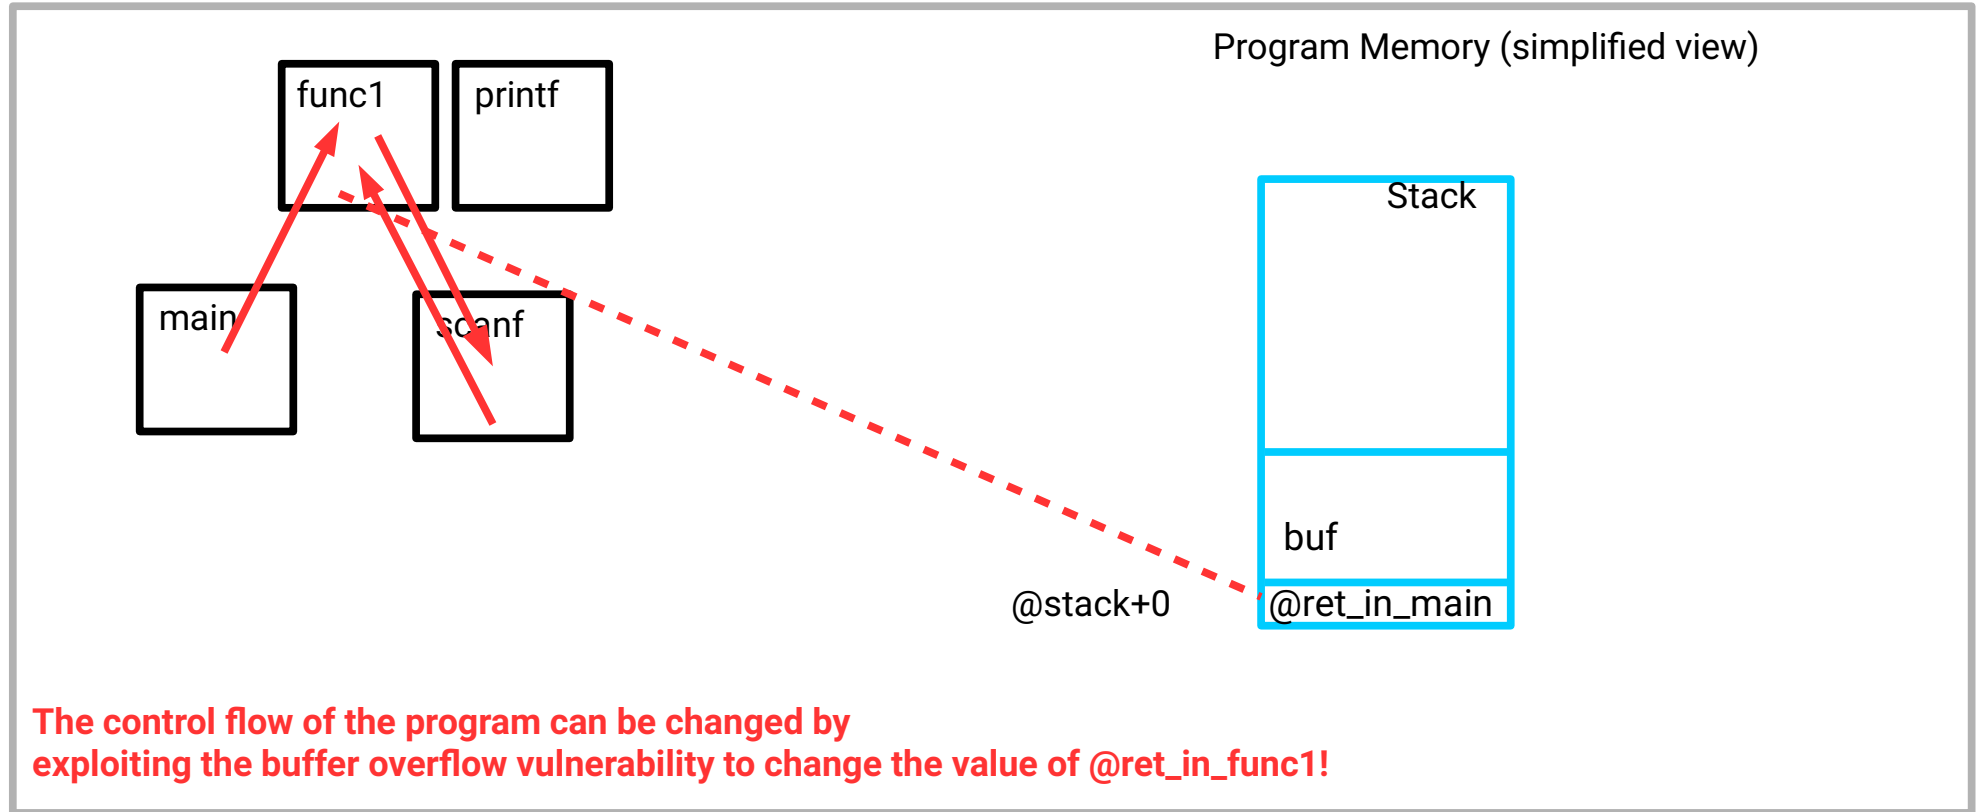
# Custom Code Execution



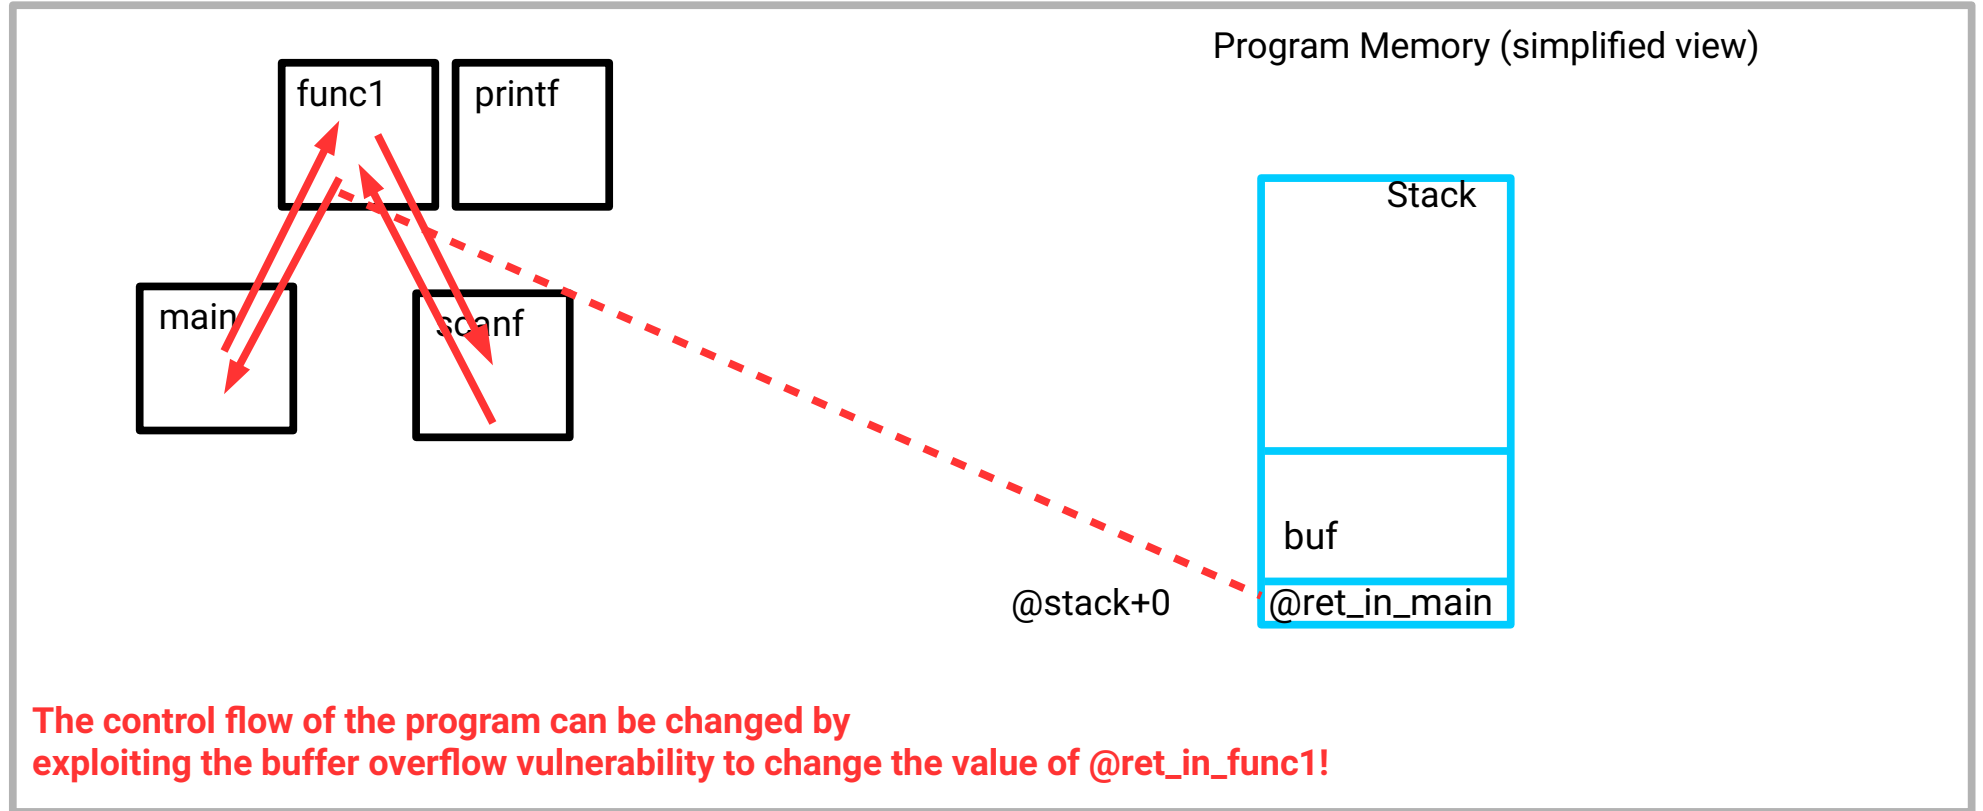
# Custom Code Execution



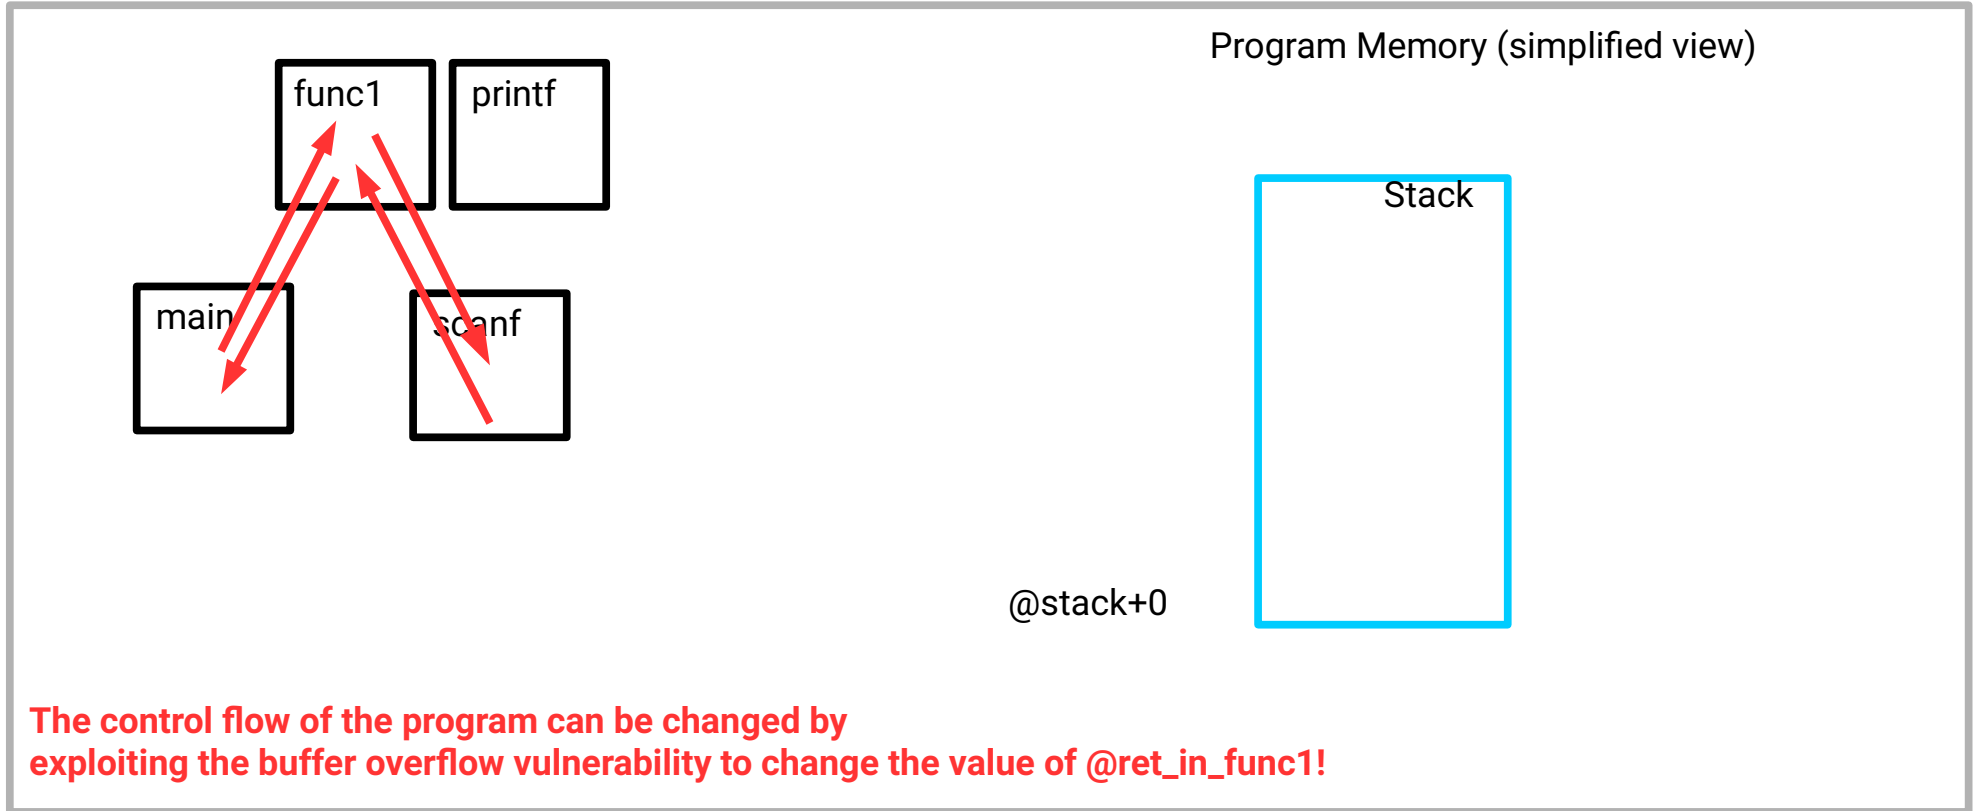
# Custom Code Execution



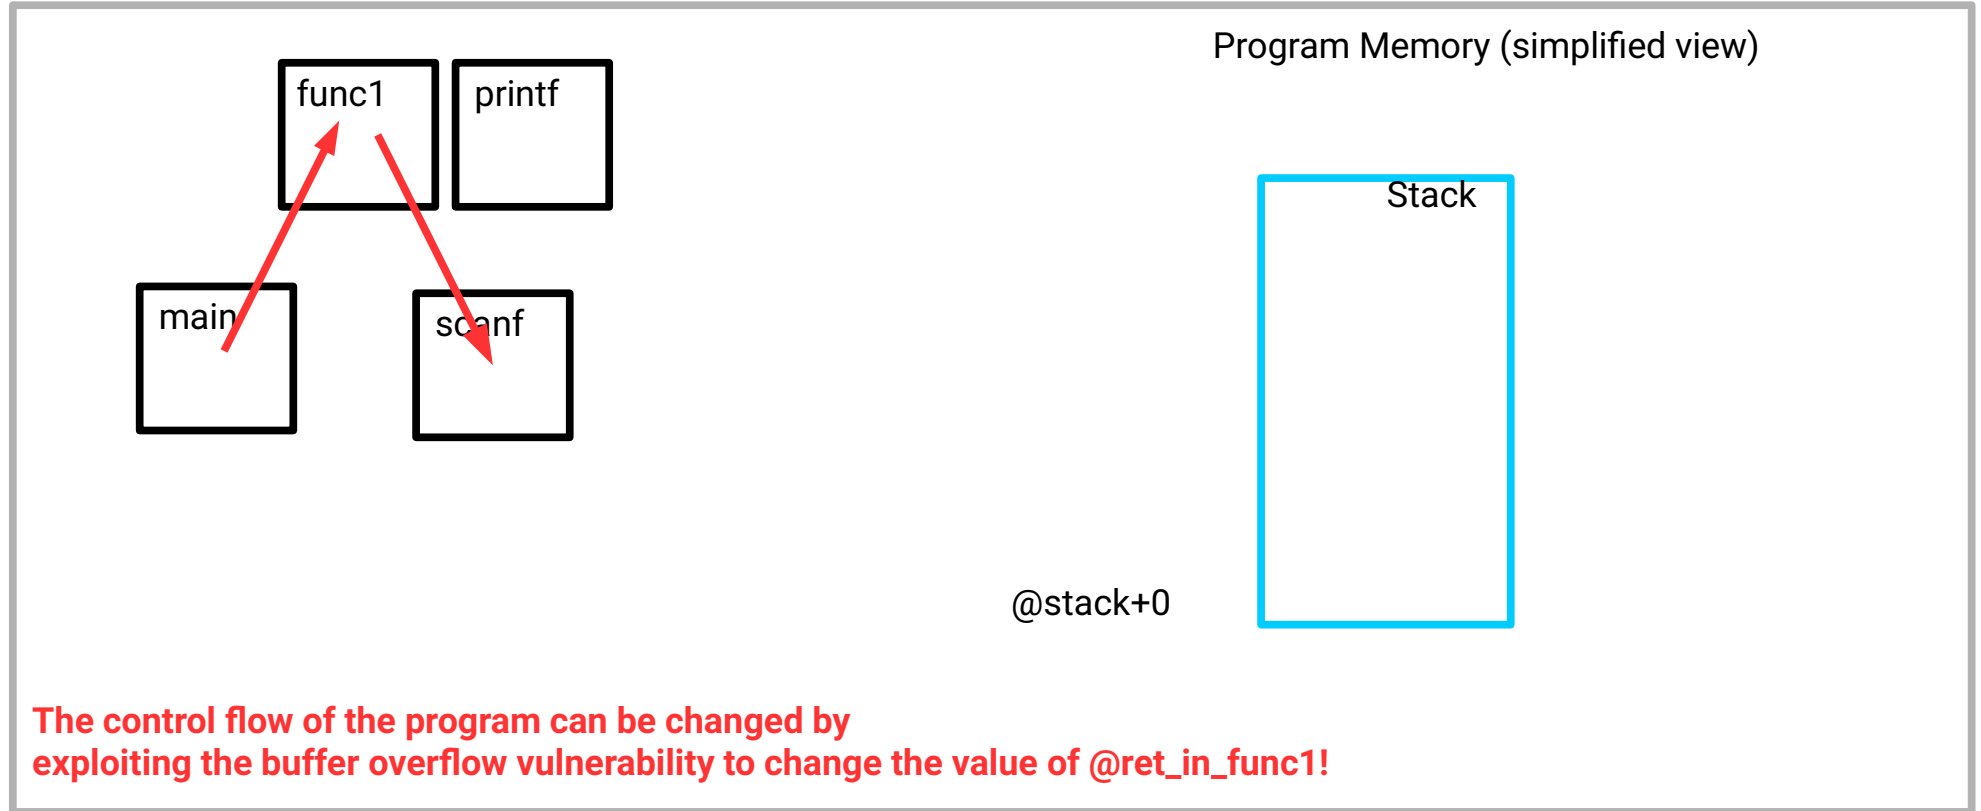
# Custom Code Execution



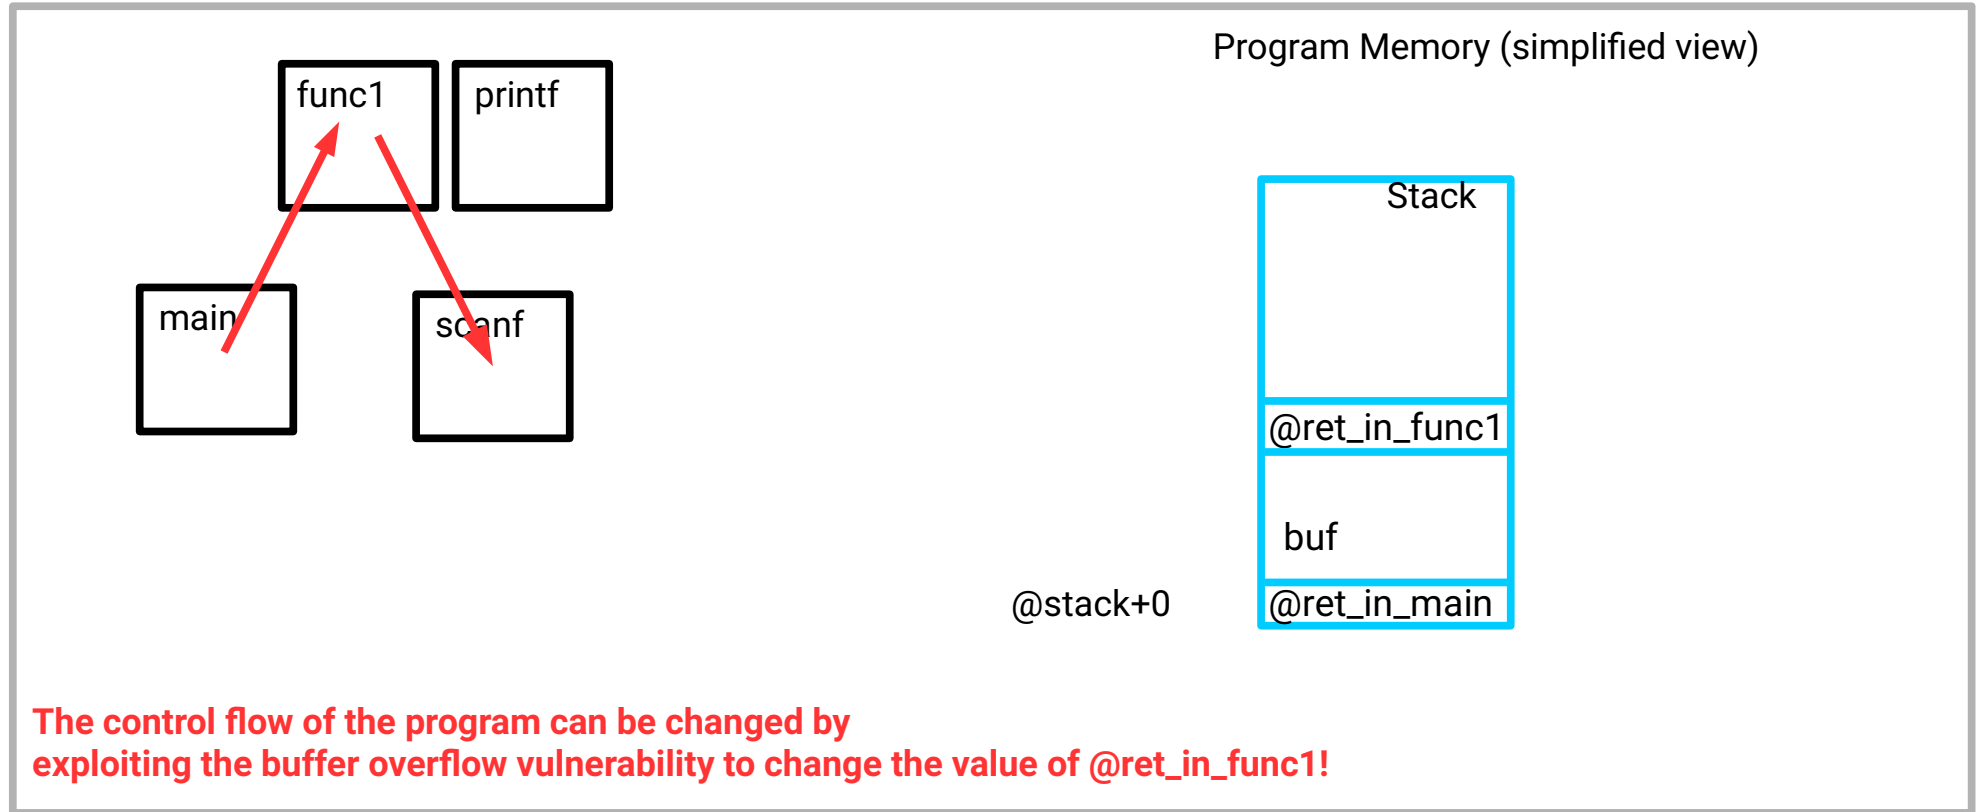
# Custom Code Execution



# Custom Code Execution



# Custom Code Execution



# Custom Code Execution



The control flow of the program can be changed by exploiting the buffer overflow vulnerability to change the value of `@ret_in_func1`!



# Custom Code Execution




The control flow of the program can be changed by exploiting the buffer overflow vulnerability to change the value of `@ret_in_func1`!

# Custom Code Execution



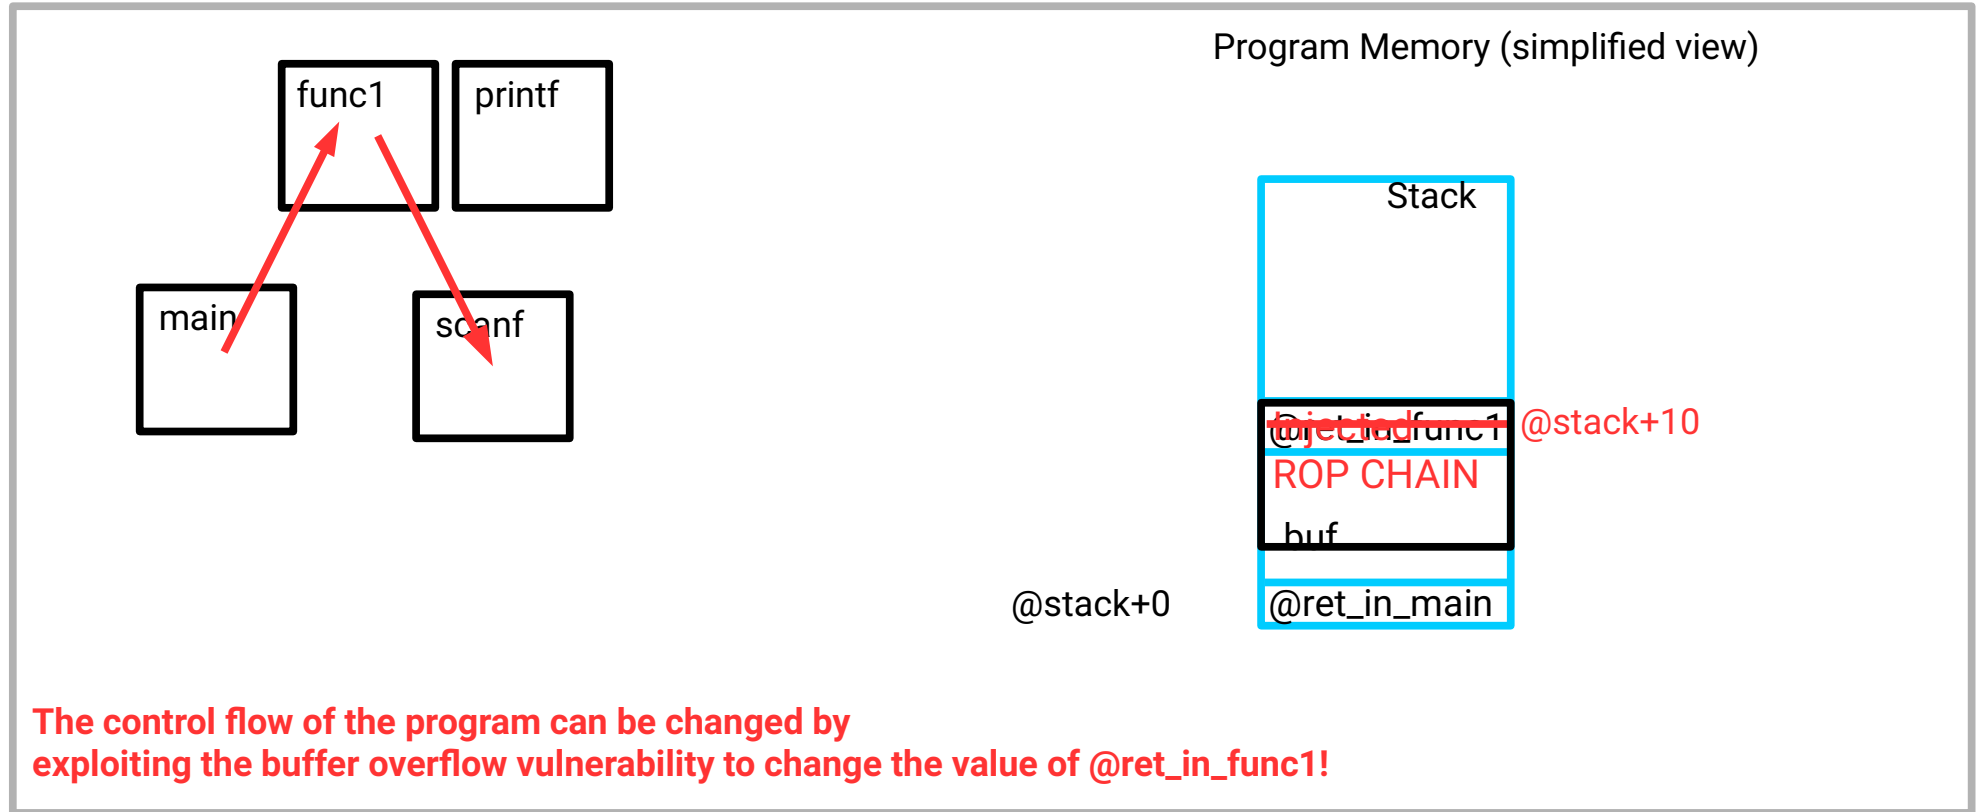
The control flow of the program can be changed by exploiting the buffer overflow vulnerability to change the value of `@ret_in_func1`!

# Custom Code Execution

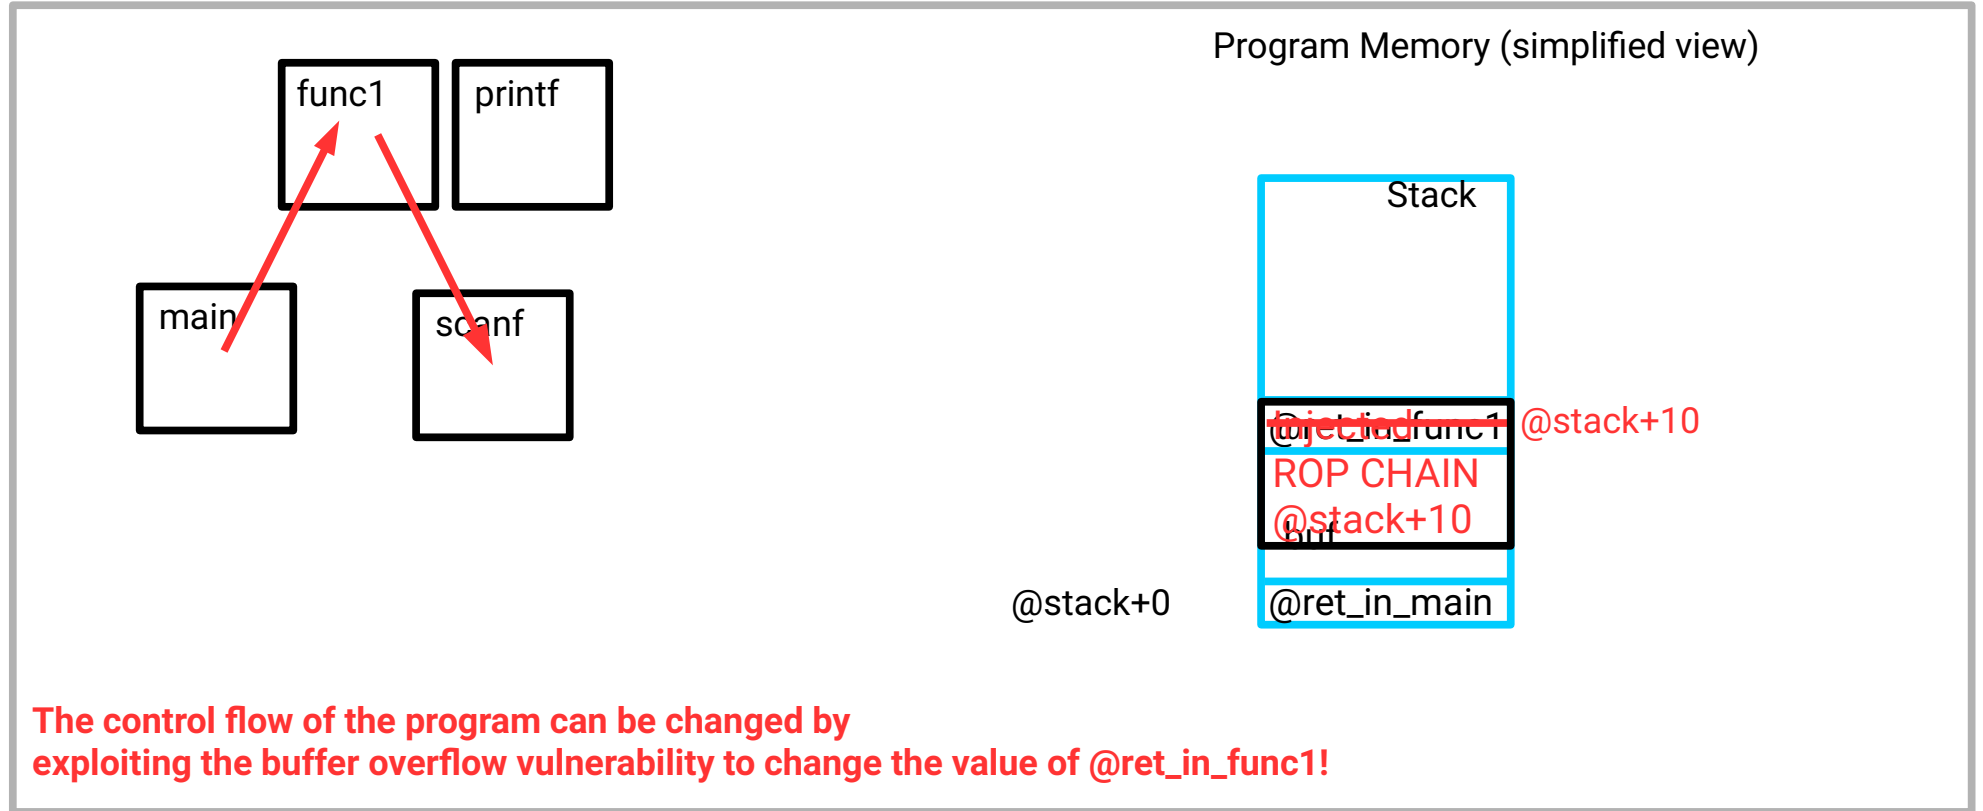


The control flow of the program can be changed by exploiting the buffer overflow vulnerability to change the value of `@ret_in_func1`!

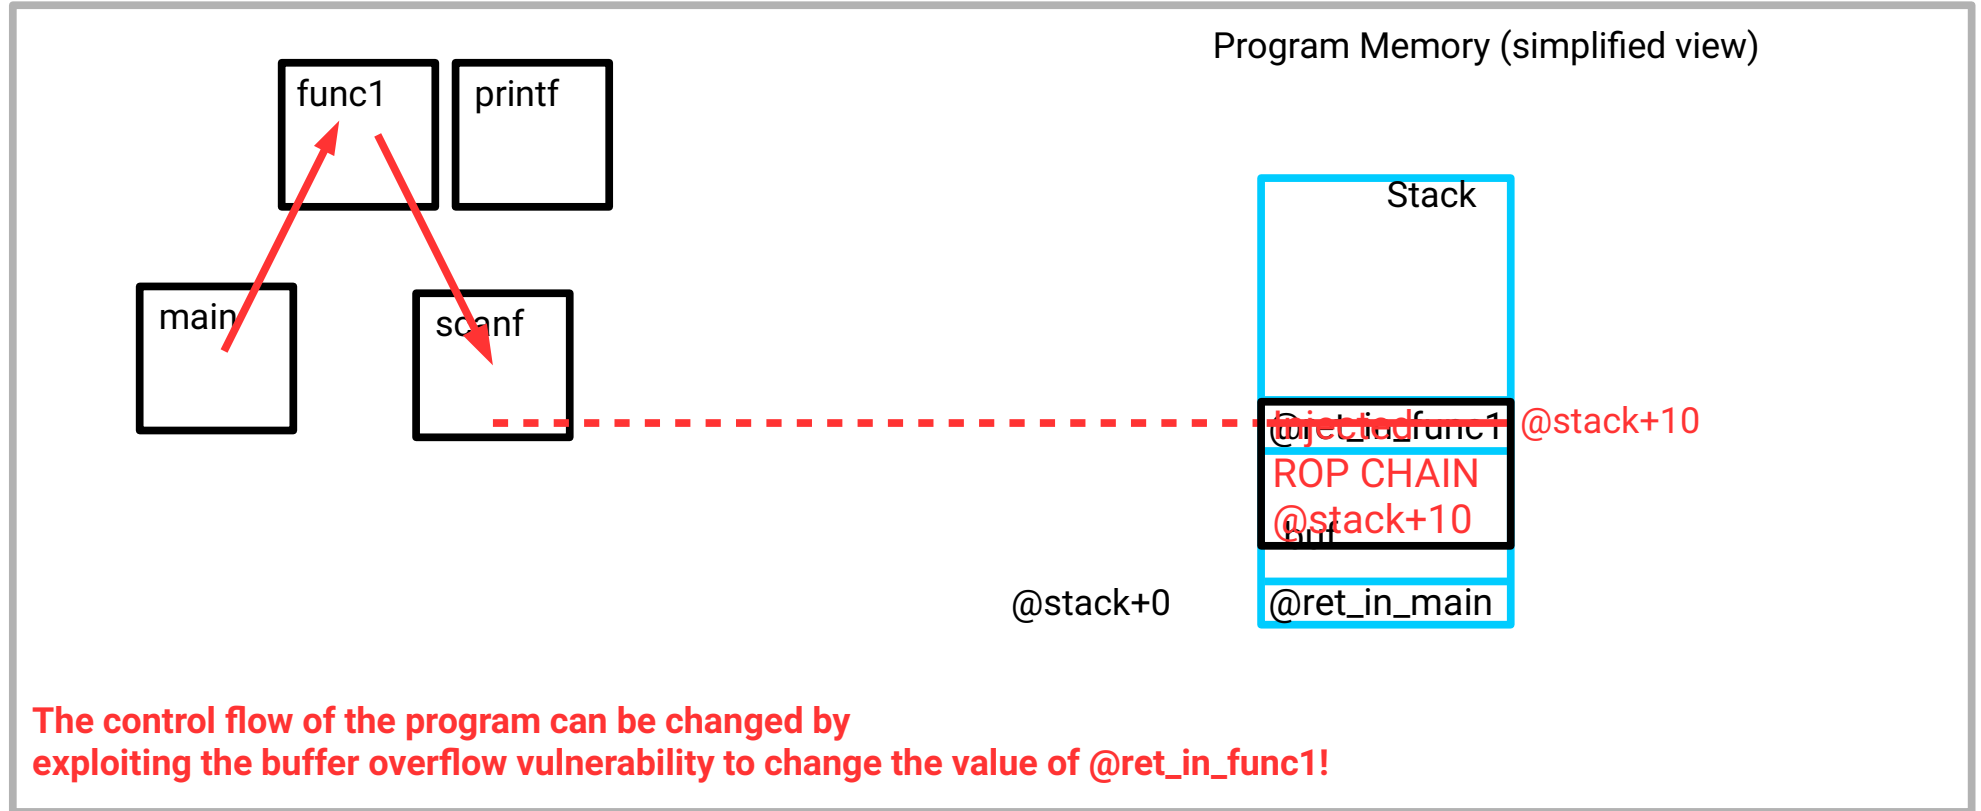
# Custom Code Execution



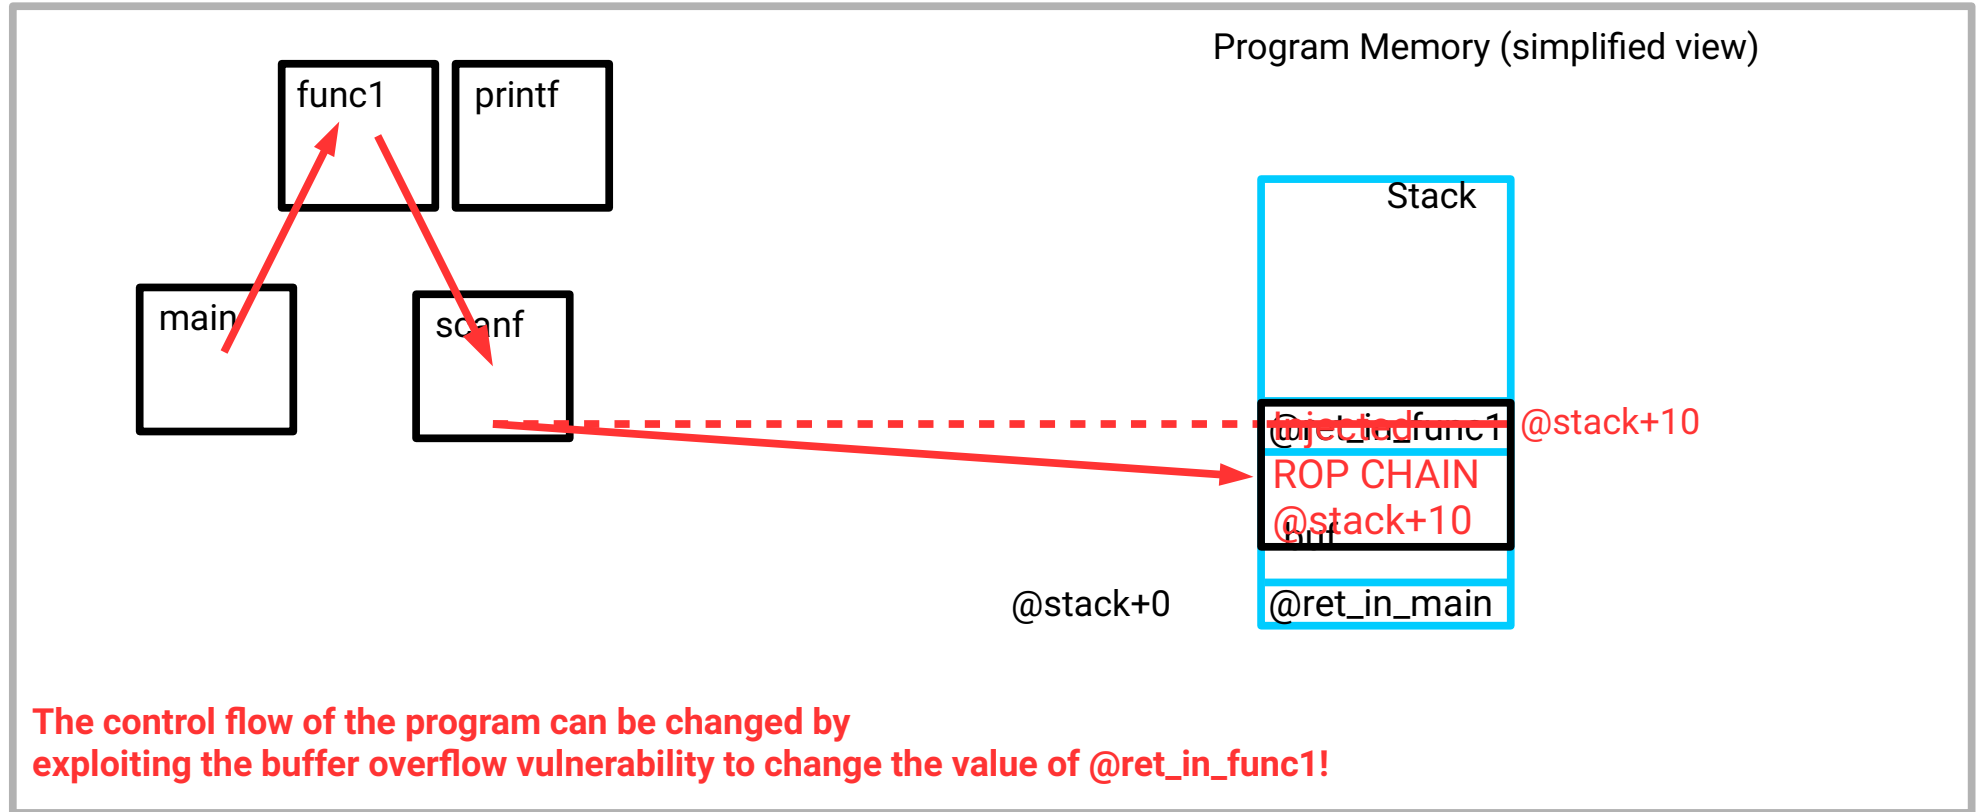
# Custom Code Execution



# Custom Code Execution



# Custom Code Execution



# Where are the Gadgets?

- Gadgets are in the code segments (program itself or linked libraries) loaded in the virtual memory of the process

```
$ cat /proc/self/maps
558a5d145000-558a5d147000 r--p 00000000 fe:01 524608 /bin/cat
558a5d147000-558a5d14c000 r-xp 00002000 fe:01 524608 /bin/cat
558a5d14c000-558a5d14e000 r--p 00007000 fe:01 524608 /bin/cat
558a5d14f000-558a5d150000 r--p 00009000 fe:01 524608 /bin/cat
558a5d150000-558a5d151000 rw-p 0000a000 fe:01 524608 /bin/cat
558a5ec90000-558a5ecb1000 rw-p 00000000 00:00 0 [heap]
7f60366d0000-7f60368f2000 r--p 00000000 fe:01 278299 /usr/lib/locale/locale-archive
7f60368f2000-7f6036914000 r--p 00000000 fe:01 403316 /lib/x86_64-linux-gnu/libc-2.28.so
7f6036914000-7f6036a5c000 r-xp 00022000 fe:01 403316 /lib/x86_64-linux-gnu/libc-2.28.so
7f6036a5c000-7f6036aa8000 r--p 0016a000 fe:01 403316 /lib/x86_64-linux-gnu/libc-2.28.so
7f6036aa8000-7f6036aa9000 ---p 001b6000 fe:01 403316 /lib/x86_64-linux-gnu/libc-2.28.so
7f6036aa9000-7f6036aad000 r--p 001b6000 fe:01 403316 /lib/x86_64-linux-gnu/libc-2.28.so
7f6036aad000-7f6036aaf000 rw-p 001ba000 fe:01 403316 /lib/x86_64-linux-gnu/libc-2.28.so
7f6036aaf000-7f6036ab3000 rw-p 00000000 00:00 0
7f6036ab3000-7f6036ab5000 rw-p 00000000 00:00 0
7f6036aed000-7f6036b0f000 rw-p 00000000 00:00 0
7f6036b0f000-7f6036b10000 r--p 00000000 fe:01 394494 /lib/x86_64-linux-gnu/ld-2.28.so
7f6036b10000-7f6036b2e000 r-xp 00001000 fe:01 394494 /lib/x86_64-linux-gnu/ld-2.28.so
7f6036b2e000-7f6036b36000 r--p 0001f000 fe:01 394494 /lib/x86_64-linux-gnu/ld-2.28.so
7f6036b36000-7f6036b37000 r--p 00026000 fe:01 394494 /lib/x86_64-linux-gnu/ld-2.28.so
7f6036b37000-7f6036b38000 rw-p 00027000 fe:01 394494 /lib/x86_64-linux-gnu/ld-2.28.so
7f6036b38000-7f6036b39000 rw-p 00000000 00:00 0
7ffd20e58000-7ffd20e79000 rw-p 00000000 00:00 0 [stack]
7ffd20f38000-7ffd20f3b000 r--p 00000000 00:00 0 [vvar]
7ffd20f3b000-7ffd20f3d000 r-xp 00000000 00:00 0 [vdso]
```



# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the *\*stack\**.  
Thus, after every pop instruction, the stack pointer, *rsp*, is moved to the next stack memory location.

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
@stack + 0x07	
@stack + 0x06	
@stack + 0x05	
@stack + 0x04	
@stack + 0x03	
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

stack grows this way ↑

# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the *\*stack\**.  
Thus, after every pop instruction, the stack pointer, *rsp*, is moved to the next stack memory location.

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
@stack + 0x07	@gadget1
@stack + 0x06	
@stack + 0x05	
@stack + 0x04	
@stack + 0x03	
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

stack grows this way ↑

# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the *\*stack\**.  
Thus, after every pop instruction, the stack pointer, *rsp*, is moved to the next stack memory location.

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
@stack + 0x07	@gadget1
@stack + 0x06	data1
@stack + 0x05	
@stack + 0x04	
@stack + 0x03	
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

stack grows this way ↑

# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the *\*stack\**.  
Thus, after every pop instruction, the stack pointer,  
rsp, is moved to the next stack memory location.

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
@stack + 0x07	@gadget1
@stack + 0x06	data1
@stack + 0x05	@gadget2
@stack + 0x04	
@stack + 0x03	
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

stack grows this way ↑

# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the *\*stack\**.  
Thus, after every pop instruction, the stack pointer,  
rsp, is moved to the next stack memory location.

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
@stack + 0x07	@gadget1
@stack + 0x06	data1
@stack + 0x05	@gadget2
@stack + 0x04	data2
@stack + 0x03	
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

stack grows this way ↑

# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the *\*stack\**.  
Thus, after every pop instruction, the stack pointer, *rsp*, is moved to the next stack memory location.

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
@stack + 0x07	@gadget1
@stack + 0x06	data1
@stack + 0x05	@gadget2
@stack + 0x04	data2
@stack + 0x03	@gadget3
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

stack grows this way ↑

# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the \*stack\*.  
Thus, after every pop instruction, the stack pointer,  
rsp, is moved to the next stack memory location.

@gadget1

Pop rax ret
----------------

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
@stack + 0x07	@gadget1
@stack + 0x06	data1
@stack + 0x05	@gadget2
@stack + 0x04	data2
@stack + 0x03	@gadget3
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

↑  
stack grows this way

# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the *\*stack\**.  
Thus, after every pop instruction, the stack pointer, *rsp*, is moved to the next stack memory location.

@gadget1 Pop rax  
ret

@gadget2 Pop rcx  
ret

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
@stack + 0x07	@gadget1
@stack + 0x06	data1
@stack + 0x05	@gadget2
@stack + 0x04	data2
@stack + 0x03	@gadget3
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

↑  
stack grows this way



# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the \*stack\*.  
Thus, after every pop instruction, the stack pointer, `rsp`, is moved to the next stack memory location.

@gadget1 Pop rax  
ret

@gadget2 Pop rcx  
ret

@gadget3 Syscall  
...

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
@stack + 0x07	@gadget1
@stack + 0x06	data1
@stack + 0x05	@gadget2
@stack + 0x04	data2
@stack + 0x03	@gadget3
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

stack grows this way



# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the *\*stack\**.  
Thus, after every pop instruction, the stack pointer, *rsp*, is moved to the next stack memory location.

@gadget1 Pop rax  
ret

@gadget2 Pop rcx  
ret

@gadget3 Syscall  
...

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
<span style="border: 2px solid red;">@stack + 0x07</span>	@gadget1
@stack + 0x06	data1
@stack + 0x05	@gadget2
@stack + 0x04	data2
@stack + 0x03	@gadget3
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

stack grows this way



# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the \*stack\*.  
Thus, after every pop instruction, the stack pointer,  
rsp, is moved to the next stack memory location.

@gadget1 Pop rax  
ret

@gadget2 Pop rcx  
ret

@gadget3 Syscall  
...

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
<span style="border: 2px solid red;">@stack + 0x07</span>	@gadget1
@stack + 0x06	data1
@stack + 0x05	@gadget2
@stack + 0x04	data2
@stack + 0x03	@gadget3
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

↑  
stack grows this way

# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the *\*stack\**.  
Thus, after every pop instruction, the stack pointer, *rsp*, is moved to the next stack memory location.

@gadget1 Pop rax  
ret

@gadget2 Pop rcx  
ret

@gadget3 Syscall  
...

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
<span style="border: 2px solid red;">@stack + 0x07</span>	@gadget1
@stack + 0x06	data1
<span style="border: 2px solid red;">@stack + 0x05</span>	@gadget2
@stack + 0x04	data2
@stack + 0x03	@gadget3
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

↑  
stack grows this way

# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the *\*stack\**.  
Thus, after every pop instruction, the stack pointer, *rsp*, is moved to the next stack memory location.

@gadget1 

Pop rax
ret

@gadget2 

Pop rcx
ret

@gadget3 

Syscall
...

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
@stack + 0x07	@gadget1
@stack + 0x06	data1
@stack + 0x05	@gadget2
@stack + 0x04	data2
@stack + 0x03	@gadget3
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

↑  
stack grows this way

# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the *\*stack\**.  
Thus, after every pop instruction, the stack pointer, *rsp*, is moved to the next stack memory location.

@gadget1 

Pop rax  
ret

@gadget2 

Pop rcx  
ret

@gadget3 

Syscall  
...

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
@stack + 0x07	@gadget1
@stack + 0x06	data1
@stack + 0x05	@gadget2
@stack + 0x04	data2
@stack + 0x03	@gadget3
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

stack grows this way



# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the *\*stack\**.  
Thus, after every pop instruction, the stack pointer, *rsp*, is moved to the next stack memory location.

@gadget1 Pop rax  
ret

@gadget2 Pop rcx  
ret

This technique is  
@gadget3 Syscall  
...

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
<span style="border: 2px solid red;">@stack + 0x07</span>	@gadget1
@stack + 0x06	data1
@stack + 0x05	@gadget2
@stack + 0x04	data2
@stack + 0x03	@gadget3
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

stack grows this way

# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the *\*stack\**.  
Thus, after every pop instruction, the stack pointer, *rsp*, is moved to the next stack memory location.

@gadget1 Pop rax  
ret

@gadget2 Pop rcx  
ret

This technique is  
Used to initialize @gadget3 Syscall  
...

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
<span style="border: 2px solid red;">@stack + 0x07</span>	@gadget1
@stack + 0x06	data1
@stack + 0x05	@gadget2
@stack + 0x04	data2
@stack + 0x03	@gadget3
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

stack grows this way



# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the *\*stack\**.  
Thus, after every pop instruction, the stack pointer, *rsp*, is moved to the next stack memory location.

@gadget1  
Pop rax  
ret

@gadget2  
Pop rcx  
ret

This technique is  
Used to initialize  
Registers with

@gadget3  
Syscall  
...

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
@stack + 0x07	@gadget1
@stack + 0x06	data1
@stack + 0x05	@gadget2
@stack + 0x04	data2
@stack + 0x03	@gadget3
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

stack grows this way

# Custom Code Execution

If the analyst inputs the following data:

'f'	'o'	@g3	'd2	@g2	d1	@g1	i0	i1	i2	...
-----	-----	-----	-----	-----	----	-----	----	----	----	-----

The pop instruction get a value from the *\*stack\**.  
Thus, after every pop instruction, the stack pointer, *rsp*, is moved to the next stack memory location.

@gadget1 

Pop rax  
ret

@gadget2 

Pop rcx  
ret

This technique is  
Used to initialize  
Registers with  
Attacker controlled values

Syscall  
...

Memory @	Value
	...
	...
...	...
@stack + 0x08	...
@stack + 0x07	@gadget1
@stack + 0x06	data1
@stack + 0x05	@gadget2
@stack + 0x04	data2
@stack + 0x03	@gadget3
@stack + 0x02	'o'
@stack + 0x01	'f'
@stack + 0x00	@code + 0x01

↑  
stack grows this way

# Finding Gadgets

'58' is the machine  
Representation  
Of "pop rax".

000000000000e76a0 <\_\_getauxval@@GLIBC\_2.16>:

[...]

e76f8: 48 8b 40 58 mov 0x58(%rax),%rax

e76fc: c3 retq

e76fd: 0f 1f 00 nopl (%rax)

e7700: 48 8b 80 10 01 00 00 mov 0x110(%rax),%rax

e7707: c3 retq

e7708: 0f 1f 84 00 00 00 00 nopl 0x0(%rax,%rax,1)

The gadget starts at address 0xe76fa, ends at address 0xe76fc and contains 2 instructions '58' (pop rax) and 'c3' (ret). The code is extracted from libc (and yes, it is possible to jump in the middle of "normal" instructions to execute "new" instructions)

# Example



# Consequences

## 1. The analyst can still put data on the stack!

- Data
- Addresses to gadgets

## 2. The injected gadgets

- Are called one after the other (they end in 'ret')
- Can put data in registers (with 'pop' instructions)
- Can call a syscall (with the 'syscall' instruction)
- "Simulate" the execution of a shellcode

# Defenses

- Two main approaches to limit exploitation of a buffer overflow
  - Stack canary (aka stack cookie)
  - Data Execution Prevention (DEP)
- Bypassing DEP
  - ROP attack with gadget chain
- Two main approaches to prevent ROP attacks bypassing DEP:
  - Address Space Layout Randomization (ASLR)
  - Control Flow Integrity (CFI)

# ROP Attacks: A bit of History

- 1997: return into-libc [1]
- 2004: attack techniques [2]
- 2007: academic description of gadgets [3]

[1] Getting around non-executable stack (and fix) - Solar Designer, 1997

[2] Pincus, Jonathan, and Brandon Baker. "Beyond stack smashing: Recent advances in exploiting buffer overruns." IEEE Security & Privacy 2.4 (2004): 20-27.

[3] Shacham, Hovav. "The geometry of innocent flesh on the bone: return-into-libc without function calls (on the x86)." ACM conference on Computer and communications security. 2007.

# Questions?

- Of course you have some questions